

Can Trust in Public Institutions Be Restored?: Evidence from Two Experimental Studies *

Guglielmo Briscese[†] Maddalena Grignani[‡]

Trust in public institutions has been declining across Western democracies, with important implications for compliance with policies that mitigate negative externalities. We test whether trust can be restored using two theory-driven experiments with representative samples of U.S. citizens. Lowering the costs of searching institution-sourced information raises factual knowledge, but improves trust only when the information is relevant to the decision-maker, with persistent effects. We discuss implications for public agencies communication campaigns.

Keywords: Institutional Trust, Inattention, Information Acquisition, Experiments.
JEL Classification Codes: D83, H83, C93

*We are grateful to the Rustandy Center for Social Sector Innovation for providing research funding for Experiment I and the TESS initiative for allowing us to field Experiment II. We thank seminar participants of the University of Chicago Experimental Economics and the University of Pompeu Fabra. We thank Marianne Bertrand, Alessandra Cassar, James Druckman, Libertad González, Ruben Enikolopov, Nicola Lacetera, and Mario Macis for their valuable comments. All errors are our own. This study obtained ethics approval from the University of Chicago IRB No. 20-0471. AEA pre-analysis plan: AEARCTR-0014160

[†]The University of Chicago gubri@uchicago.edu

[‡]Pompeu Fabra University maddalena.grignani@upf.edu

1 Introduction

Trust in public institutions is fundamental to economic stability and social cohesion in liberal democracies. Since Adam Smith, scholars have argued that when citizens view institutions as competent and trustworthy, they are more likely to comply with recommended behaviors and believe that others will do the same, creating a virtuous cycle that enhances social welfare, from taxation to public health (Besley and Dray, 2024; Levi and Stoker, 2000; Tyler, 1990; Levi, 1988; Smith, 1759, 1776). Consistent with this correlation, higher institutional trust is associated with greater compliance with taxation and public health guidance, such as vaccination (Tyler, 1990; Levi and Stoker, 2000; Darden and Macis, 2024; Besley and Dray, 2022; Fancourt et al., 2020; Martinez-Bravo and Stegmann, 2022). Yet trust in public institutions has been eroding across much of the Western world, with particularly steep declines in the U.S.: in 2024, only 37% of Americans reported trusting public institutions, compared to 70% in 1972 (Gallup, 2024). This erosion has been especially pronounced along partisan lines, with trust in specific institutions closely tracking political affiliation and exposure to partisan media narratives.

A large body of literature has documented that institutional trust can erode quickly when citizens encounter information – accurate or not – that questions an institution’s integrity or performance: institutional failures, poor communication choices, elite political attacks, and partisan media all play a role in reducing trust in institutions and willingness to comply with recommended behaviors, even when non-compliance carries substantial personal and societal costs (Feldman et al., 2012; Hmielowski et al., 2014; Algan et al., 2021; Suhay et al., 2022; León-Ciliotta et al., 2023; Böhm et al., 2026; Darden and Macis, 2024; Martinez-Bravo and Stegmann, 2022; Alsan and Wanamaker, 2018). Evidence on what scalable interventions can help *rebuild* trust is scarcer. A smaller set of studies has shown, for example, that educating citizens about a central bank’s independence raises trust in the institution (Brouwer and De Haan, 2026), and communicating the bank’s institutional structure and objectives mitigates partisan misperceptions of it (Kuang et al., 2025). Similarly, exposing voters to the procedural safeguards used by election officials in other states raises confidence in elections among Republicans, Democrats, and Independents alike (Gaudette et al., 2026). This evidence suggests that exposure to belief-correcting information can improve trust, but it is limited and presents some methodological limitations. For example, effects are often measured immediately after the intervention, so little is known about whether they persist and to what extent social desirability bias may partially drive them. Moreover, the information is shown unconditionally to a randomized set of respondents rather than sought by them, leaving open whether citizens would acquire it on their own if given the opportunity-cost outside of a study setting. Further, because these designs rarely vary the content of what is provided while holding fixed that any information is provided, it is often unclear whether an observed effect reflects the value of the information itself or simply exposure to the institution. This distinction carries different policy implications: if trust responds to exposure alone, outreach campaigns can succeed by maximizing contact, while if it responds only to information that citizens consider useful, campaigns must be built on delivering relevant information, an arguably more demanding effort.

This study addresses these gaps by offering novel causal evidence and unpacking the mechanisms behind the efficacy of scalable information campaigns. We begin with a standard model of rational inattention that incorporates demand for institutional information: citizens’ attention is scarce and allocated to information that is valuable to make decisions, and ignoring an institution can be optimal when benefits are unclear. A key hypothesis behind our study is that citizens may lack familiarity with the work of public institutions and how it benefits them. While they could seek information to update their beliefs, they may be unwilling to incur the search costs if the perceived benefits are not immediately clear. In the absence of belief-correcting information, citizens rely on politically motivated heuristics, producing biased judgments about institutional trustworthiness. Our model makes the resulting trap explicit: trust updates only upon consumption, so a citizen whose trust falls below the threshold never acquires the evidence that would revise it. Mistrust persists not because it is confirmed but because it is never tested — a vicious cycle in which individuals avoid engaging with valuable institution-provided information, such as updates on public health, environmental risks, or product safety regulation, further entrenching their mistrust and reducing compliance with welfare-enhancing recommendations. From these hypotheses follow the main questions we take to the data: can lowering the search cost increase consumption of institution-sourced information and durably improve trust in the institution? And does the *value* of the information provided matter? We test these theory-driven questions in two experiments where we exogenously vary search costs and information relevance to assess whether scalable interventions can increase engagement with institutional information, and consequently enhance trust and willingness to follow institutional guidance.

In our first experiment, we recruited a representative sample of Americans to participate in a seven-wave longitudinal survey from April to October 2020. This was a period of intense political polarization and heightened skepticism toward government agencies, providing a demanding test of whether such interventions can work even under adverse conditions. We find a rapid decline in trust, particularly in institutions managing the pandemic (i.e., the scientific community and the healthcare sector). The richness of our data allows us to demonstrate how this erosion is more strongly correlated with misperceptions about the crisis and partisan media consumption than with personal negative experiences. In the sixth wave, implemented in June 2020, we provided a random half of the respondents with an optional, non-incentivized link to the online COVID-19 death tracker managed by the CDC — the Centers for Disease Control and Prevention, the institution at the core of the COVID-19 pandemic management in the US. This randomized intervention reduced the cost of accessing a clear, verifiable metric of the ongoing crisis’s severity, without forcing respondents to consume it. We measure the effect of this light-touch intervention on knowledge and trust outcomes. We also introduced several features to minimize risks of social desirability bias and experiment demand effects: (i) a four-month delay between treatment and outcomes, with the in-between period dominated by other topics on the Presidential election campaign, (ii) placebo institutions shown in random order in our outcome questions, (iii) long-term obfuscated follow up knowledge accuracy question that could not be answered using the original source, and (iv) redirected respondents’ attention on another hotly debated topic (i.e., the George Floyd protests) after the treatment to reduce salience.

We find that, four months later, the treatment increased accurate knowledge about the ongoing pandemic by 5.4 percentage points from a control-group mean of 0.53 ($p = 0.08$), the share reporting higher trust in the CDC by 6.2 percentage points from a control-group mean of 0.12 ($p = 0.03$), and perceptions of CDC political independence by 6 percentage points from a control-group mean of 0.08 ($p = 0.03$). Treated respondents were also 6.6 percentage points more likely to have taken a COVID-19 test from a control-group mean of 0.28 ($p = 0.03$), though effects on other behavioral outcomes appear to be more muted. Long-term meta-cognitive measures show that improvements in trust reflect conscious belief updating.

We implemented a second experiment four years later, in August 2024, fielded ahead of a Presidential election so that both studies test the intervention under similarly polarized political conditions. This second study focused on trust in the Environmental Protection Agency (EPA), one of the most debated public institutions during that political campaign. We recruited a representative sample of Americans and randomly assigned them to different conditions: two groups received a factual question about current air quality in their zip code, one with a direct link to the EPA’s tracker (low search cost), one without (high search cost). A third randomized group received an equivalent low-cost link for a lower-value question about the location of their EPA regional office. Holding the institutional source constant across all arms reduces the likelihood that differences in outcomes reflect messenger endorsement or experimenter demand, thereby making the effects more plausibly attributable to the content and relevance of the information itself. We find that low awareness of the EPA is correlated with low baseline trust, consistent with our rational inattention hypothesis. Reducing search costs significantly improves knowledge accuracy: close to half of respondents with a link answered correctly, compared to fewer than one in five without ($p < 0.01$). Respondents who accessed high-value air quality data were 6.2 percentage points more likely to view the EPA as politically independent than those who accessed lower-value office location information — a 34 percent increase on the office-location arm’s base of 18 percent ($p = 0.01$) — even at equivalent search costs, suggesting that information value, and not mere institutional exposure, shapes trust. Effects on other trust dimensions - benevolence (the EPA improves lives) and competence (support for the EPA’s involvement in local research activities) - are directionally consistent.

Across our two studies, we provide novel causal evidence for a mechanism linking informational frictions to institutional mistrust: citizens who face high costs of engaging with institution-sourced information may default to ideological shortcuts (Bénabou and Tirole, 2011), and targeted outreach campaigns may reduce search costs, helping break this cycle. While rebuilding trust in public institutions is likely to be a multifaceted challenge, our results suggest that large-scale digital information campaigns, often implemented by public institutions, may help achieve this goal, provided they prioritize delivering timely, relevant content rather than initiatives focused solely on increasing institutional visibility.

The paper is structured as follows. In the subsection below, we place our study into the existing literature and point to our contributions to different streams of research. Section 2 presents a model of rational inattention that motivates our hypotheses and experimental

design. Sections 3 and 4 describe the data, methodology, and results of Study I and II respectively. Section 5 concludes with suggested directions for future studies and policy implications.

Related literature Our study contributes to several streams of literature. First, we add to the economics of trust, which predominantly examines interpersonal trust as a driver of economic prosperity (Sapienza et al., 2013; Guiso et al., 2011; Fehr et al., 2002). Institutional trust has received comparatively less attention; existing work shows it can erode rapidly in response to failures or negative media coverage, reducing compliance with welfare-enhancing behaviors such as vaccination (Rafkin et al., 2021; Darden and Macis, 2024; Alsan and Wana-maker, 2018; León-Ciliotta et al., 2023), preventive health interventions (Algan et al., 2021), and tax compliance (Besley and Dray, 2024), but offers little guidance on restoration strategies. A notable exception is Acemoglu et al. (2020), who show in a lab-in-the-field experiment in Pakistan that information about positive judicial reforms raises trust in courts. Our study differs from Acemoglu et al. (2020) in three key respects: our design arguably offers higher external validity by allowing respondents to self-select into information engagement; we track effects over several months using obfuscated follow-up measures; and we use meta-cognitive questions to assess whether changes in trust are conscious.

Second, we contribute to political science research on institutional trust. Foundational theories identify competence, integrity, and legitimacy as the pillars of trust (Levi, 1988; Tyler, 1990; Mayer et al., 1995), while empirical evidence on interventions addressing biased institutional judgments remains scarce (Grimmelikhuijsen and Meijer, 2014; Van de Walle and Bouckaert, 2003). A related strand argues that distrust arises when citizens perceive misaligned incentives between public officials and citizens (Besley, 2006; Alesina and Tabellini, 2007), consistent with the strong correlation between perceived politicization and low trust (Bolsen and Druckman, 2018; McLaughlin et al., 2021; Van Scoy et al., 2021; Clark et al., 2023) and evidence that politically neutral scientists are viewed as most credible (Alabrese et al., 2024). We provide causal evidence that reducing informational frictions can improve perceived political independence, one component of trust in both frameworks.

Third, we contribute to the literature on government communication campaigns. While Banerjee et al. (2024) document that SMS reminders improved compliance with COVID-19 guidelines, other studies report more heterogeneous effects that depend on contextual features and the perceived relevance of the information (Carson-Chahhoud et al., 2017; Ravalion et al., 2015). Building on the behavioral inattention literature, which shows that limited attention can lead to suboptimal decisions (Gabaix, 2019; DellaVigna, 2009), we propose a simple framework that helps reconcile these mixed findings and provide evidence that lowering access costs to relevant information, rather than simply increasing exposure to generic institutional messaging, may be a more effective policy lever. Our results suggest that politically polarized perceptions of institutions partly reflect inattention: individuals who do not actively engage with institution-provided information may rely on ideological heuristics, and targeted outreach campaigns can partially mitigate this tendency.

Fourth, we connect this literature to work on attention and the demand for information. Rational inattention holds that attention is scarce and allocated to what matters for the decision at hand (Sims, 2003; Gabaix, 2014); consistent with this, attention to macroeconomic news rises with its payoff relevance (Link et al., 2026), inattention shapes the treatment effects recovered in randomized trials (Maćkowiak and Wiederholt, 2025), and merely making an attribute salient changes choices (Conlon, 2026). We also build on the survey-experimental literature on information and belief updating (Haaland and Roth, 2023; De Quidt et al., 2018; Fuster et al., 2022; Martinez-Bravo and Sanz, 2022), which supplies the design tools we use to elicit and move beliefs credibly. What is missing from both strands is a link between the *demand* for the information an institution produces and *trust* in that institution. Our contribution is to supply it: we model trust as expected usefulness, which enters the attention problem directly and updates only when the information is actually consumed. That single change turns a standard inattention model into one with a self-sealing trap, and it is what our two experiments are designed to test. The remainder of the paper is structured as follows. Section 2 presents the rational inattention model. Section 3 details the first experiment. Section 4 presents the second experiment. Section 5 concludes.

2 Theoretical Framework

We ground our experimental design in a model of rational inattention (Sims, 2003), using the sparsity formulation of Gabaix (2014, 2019), which yields a closed form in which attention falls to exactly zero below a threshold. Citizens are attention-constrained: they face costs when searching for or processing information, and ignoring a source altogether can be optimal. An individual allocates attention $m_i \in [0, 1]$ to each attribute x_i relevant to a decision, trading off the benefit of being better informed against the cost of acquiring and processing the information. That cost rises with a search-cost parameter $c_i > 0$, specific to each piece of information, and with an individual-level cognitive cost $\kappa > 0$.

We depart from the standard framework in one respect, which is central to our argument. We define *trust* in institution j as $\theta_j \in [0, 1]$, the share of j 's work that the citizen expects to be useful to her. This is a deliberately narrow, operational definition: it is closest to the competence dimension in Mayer et al. (1995), and it parallels the way beliefs about an information source's accuracy are modeled in previous studies (Gentzkow et al., 2025). Trust so defined enters the agent's problem directly, because an attribute the citizen expects to be useless is costless to ignore: the expected loss from inattention to an attribute produced by j is scaled by θ_j . The optimal attention devoted to attribute x_i is then

$$m_i^* = A\left(\frac{\theta_j \sigma_i^2 |a_{x_i} u_{aa} a_{x_i}|}{c_i \kappa}\right), \quad A(z) = \max\{0, 1 - 1/z\} \quad (1)$$

where σ_i^2 captures the variability of the information (how informative it is), a_{x_i} how much a change in x_i shifts the optimal action, and u_{aa} the stakes of that action. Together these terms imply that an individual pays more attention to information that is more variable, more consequential for decisions, and less costly to acquire — and that comes from an institution she expects to be useful.

Because $A(\cdot)$ is zero below a threshold, the model has a corner solution. The citizen

consults institution j about attribute i if and only if

$$\theta_j > \bar{\theta}_i \equiv \frac{\kappa c_i}{\sigma_i^2 |a_{x_i} u_{aa} a_{x_i}|}. \quad (2)$$

Low trust, high search costs, or information of little relevance to any decision can each push a citizen below the threshold and produce *zero* attention, not merely less of it. This corner is what generates a vicious cycle, because trust updates only upon consumption. A citizen who consults j observes whether the information was in fact useful and revises θ_j toward that experience; a citizen who does not consult j observes nothing and leaves θ_j unchanged. Someone whose trust starts below $\bar{\theta}_i$ therefore never acquires the evidence that might revise it, and mistrust persists whether or not it was warranted in the first place. Appendix A1.1 gives the full derivation, including the updating rule, and Appendix A1.1.8 works through a numerical example showing how the arms of our two experiments map onto the parameters of equation (1).

Our experiments intervene on the right-hand side of equation (2). Lowering the search cost c_i — the CDC link versus no link in Study I, and the link-versus-no-link contrast in Study II (arms T1 and T2, described in Section 4 — lowers the threshold. Varying the decision-relevance a_{x_i} of what lies behind the link — T1 versus T3 in Study II — changes it as well. If the intervention brings citizens over the threshold, they consult, they learn whether the information was useful, and θ_j can move. This also clarifies what the model predicts for the low-value arm. Lowering c_i for information with a small a_{x_i} still brings some citizens over the threshold — fewer than for high-value information, since the threshold is higher — but what they find there tells them less about how useful the institution’s work is, so trust moves less. T3 should therefore show positive but lower engagement than T1, and a smaller effect on trust.

One further distinction, which we use to interpret which dimension of trust moves in Study II, is worth making explicit. An attribute can fail to be useful for two reasons: because it bears on no decision the citizen takes, or because it is distorted — reported in a way that serves the institution’s political principals rather than the citizen. Perceived politicization is the citizen’s belief about the second. The two failures are learned about in different ways. Relevance is largely known from the question itself: a respondent told that the EPA tracks air pollution in real time already knows whether that bears on her decisions. Distortion is not: it is learned by consulting a number that can be checked — against the air outside, a weather application, or the local news — and finding it accurate. Consulting a verifiable, current figure therefore updates beliefs about distortion more than beliefs about relevance, and the model predicts that the T1–T3 contrast should move perceived political independence more than perceived benevolence, and should do so most for respondents in a position to verify the number. We stress that this reading is an interpretation offered after seeing which outcome moved, not a prediction the design was built to test; Remark 3 in Appendix A1.1 states the decomposition formally.

2.1 Empirical Predictions

The model yields three testable predictions that map directly onto our experimental designs.

Our first hypothesis is that lowering search costs will increase an individual’s attention. While this hypothesis is consistent with previous studies on rational inattention, it has not been investigated in a context in which individuals may hold polarized beliefs about the institution providing such information. Would reducing search costs be enough to partly offset the information avoidance tendency driven by strongly held beliefs? Our first hypothesis thus reads as follows:

H1: Reducing search costs (c_i) increases attention (m_i^).*

Our second hypothesis is that increased attention leads to improved knowledge accuracy. This hypothesis also deserves to be tested empirically since it is unclear a priori whether citizens do in fact revise their factual knowledge as a result of consuming information from a public institution, or whether they discard it if not in line with their prior beliefs:

H2: Reducing search costs (c_i) increases knowledge accuracy.

Our third hypothesis is that after consuming institutional information, individuals will appreciate its usefulness for their decision-making and update their beliefs about institutional capabilities, thereby increasing their trust.

H3: Consuming high-value information increases trust in the institution providing the information.

Finally, the model focuses on individual-level attention costs and information value as determinants of trust. It abstracts from the macro-level forces behind the sharp aggregate decline in institutional trust we document in Section 3.2, including elite political attacks, partisan media narratives, and the pandemic shock. We view these as complementary rather than competing mechanisms: the former help explain why trust declined so sharply, while our framework motivates why it may be partly recoverable through targeted informational interventions. Our experiments test this more focused question: conditional on these political headwinds, can reducing attention costs still shift trust?

3 Study I: Search Costs, Information Acquisition, and Trust Beliefs

3.1 Data

Our first study focuses on Americans’ trust in the U.S. Centers for Disease Control and Prevention (CDC). During the COVID-19 pandemic, the CDC became a central public institution, guiding the national response with health recommendations and policies. However,

its role was frequently debated along political lines, with conflicting opinions on its guidance on lockdowns, mask mandates, and vaccines (see selected media extracts in Figure A2 in the Appendix). We partnered with NORC, the organization that administers the General Social Survey (GSS), using their AmeriSpeak Panel.¹ We recruited 1,440 U.S. citizens from the AmeriSpeak panel to participate in a seven-wave longitudinal survey conducted between April and October 2020 (see Table A8 in the Appendix for demographic and socioeconomic characteristics).² This longitudinal design allowed us to track shifts in beliefs over time, observe changes in the absence of an intervention for the control group, evaluate the long-term effects of our experimental intervention, and reduce the risk of social desirability bias. The first wave, conducted in early April 2020 as COVID-19 deaths surged across the U.S.,³ gathered baseline data on trust in institutions, media consumption, and beliefs, as well as information on respondents’ personal experiences with the crisis. The following three weekly waves tracked lived experiences during the critical first month of the pandemic. Wave 5, conducted the week of May 18, measured perceptions of the crisis’s severity. Wave 6, administered on June 22, 2020, contained the randomized information provision intervention. The final wave, conducted the week of October 19, 2020, was timed to capture changes in beliefs and preferences four months after the intervention and just before the presidential election. While respondents were incentivized to complete all waves, 1,076 of the original 1,440 took part in the final wave, and 990 of these had also taken part in wave 6, when treatment was assigned; the October outcome regressions are estimated on these 990 respondents.⁴ Attrition is correlated with several demographic characteristics: income, age, Hispanic ethnicity, cohabitation status, parenthood, and direct COVID-19 exposure. Importantly, attrition does not differ by political beliefs, news consumption, or treatment assignment (see Table A9 in the Appendix). We address selective attrition through NORC-provided survey weights, calibrated to population benchmarks to maintain representativeness across key demographic dimensions, which we use throughout our analyses. We further examine whether treatment effects vary by the characteristics that predict attrition to assess whether selective retention biases our estimates.

3.2 Pre-experimental findings

We report here the descriptive results from the first five waves of our longitudinal survey, which motivated the experimental intervention in the sixth wave. The most striking pattern

¹Funded and operated by NORC at the University of Chicago, AmeriSpeak[®] is a probability-based multi-client household panel designed to be representative of the U.S. household population, covering approximately 97% of households. NORC places particular emphasis on recruiting underrepresented groups, including households without internet access, providing them with connectivity and tablets to enable participation.

²This was part of a larger study on COVID-19 that also collected data on respondents’ mental health and well-being — topics unrelated to the present study — gathered in a way that did not interfere with our experimental design (Bertrand et al., 2020).

³See Figure A3 in the Appendix for the evolution of the COVID-19 death rate alongside our survey waves.

⁴The 86 wave-7 respondents who skipped wave 6 were never assigned to a treatment arm and are therefore excluded from the experimental analysis. Item non-response accounts for the remaining differences in sample size across outcomes (990 for perceived independence, 978 for the self-reported change in confidence, 977 for the knowledge ranking).

is the rapid erosion of trust in public health institutions — specifically hospitals, healthcare professionals, and the scientific community — over just four weeks, with a similar decline observed for health insurance companies. Our repeated survey waves allowed us to track this erosion closely and confirm that trust did not return to pre-pandemic levels by the final wave (see Figure A4 in the Appendix). This decline in trust was closely linked to compliance. Table A10 in the Appendix shows a strong correlation between trust in health-care institutions and self-reported adoption of preventive behaviors, both in levels and in trends. Respondents with higher trust in health institutions were more likely to follow recommended guidelines — handwashing, physical distancing, and mask-wearing — and those who increased their trust between the first and last week of April were also more likely to report adopting additional preventive behaviors over the same period. Notably, we do not observe this correlation with trust in government more broadly (the White House and the Federal Government), consistent with quasi-experimental evidence that declining trust in specific institutions leads to suboptimal health behaviors (Martinez-Bravo and Stegmann, 2019). Partisan media consumption amplified these differences: when asked to project the number of additional COVID-19 deaths by year’s end, respondents consuming left-leaning news estimated a figure twice as high as those consuming right-leaning news (Figure A5 and Table A12 in the Appendix). Taken together, these patterns document a rapid, politically polarized erosion of trust — exacerbated by slanted media consumption — that was associated with lower self-reported compliance with behaviors carrying significant negative externalities. This motivates our core research question: can scalable interventions improve individuals’ understanding of the crisis, rebuild institutional trust, and increase their willingness to follow public health recommendations?

3.3 Experimental methodology

Our experiment tests whether mistrust in public institutions stems partly from citizens’ limited familiarity with institutional work — and their unwillingness to incur the search costs required to remedy this. The experiment reduces these costs and allows individuals to self-select into consuming institution-sourced information, simulating real-world settings where credible information is made readily accessible through public information campaigns or timely crisis communications. Respondents were asked a verifiable question about the COVID-19 death rate — a quantifiable, personally relevant measure of the crisis. Half were randomly assigned to a treatment group and offered a prompt encouraging them to look up the answer before responding: *“If you wish to do so, you can look up the answer on the official CDC website at the following link: <https://www.cdc.gov/coronavirus/2019-ncov/cases-updates/cases-in-us.html>”*. Respondents were first asked to estimate the coronavirus death toll in their state on an eight-point categorical scale,⁵ and then to estimate the nationwide death toll on a sliding scale from 0 to 200,000 in increments of 20,000.⁶ The nationwide question was deliberately framed differently from the state-level question to minimize anchoring.

⁵*“How many people have died in your state because of coronavirus from the first death until today?”*
Response options: *Less than 500; Between 500 and 1,000; Between 1,000 and 3,000; Between 3,000 and 5,000; Between 5,000 and 10,000; Between 10,000 and 15,000; Between 15,000 and 30,000; More than 30,000.*

⁶*“How many people have died in the U.S. because of coronavirus from the first death until today?”*

Both questions appeared in the first block of ten questions in that wave to reduce survey fatigue. The link was spelled out in full so that respondents who distrusted the CDC could choose not to click it. Those who did were redirected in a new browser window to the CDC’s national statistics page, where an interactive map allowed them to view state-level death counts by hovering over their state. We matched each respondent’s answers to official CDC statistics from the day of survey completion to assess accuracy. These were the only questions about the current COVID-19 death rate across all seven waves, and the only occasion on which treated respondents were exposed to CDC-sourced information. Randomization was performed using the Stata package *randtreat*,⁷ stratifying on gender, age, education, urban/rural residence, and political affiliation. This design allowed participants to voluntarily engage with the CDC — motivated by the opportunity to verify their own knowledge — without coercion and without financial incentives, preserving the naturalistic character of information acquisition.^{8 9}

3.3.1 Addressing Experimenter Demand and Messenger Effects

A central concern in information-provision experiments is whether treatment effects reflect genuine belief updating or experimenter demand — respondents inferring the researcher’s objectives and adjusting answers accordingly. Khan et al. (2021) show that modest positive effects of an information campaign in Pakistan were largely attributable to experimenter demand, underscoring the importance of accounting for this mechanism. Although our design does not include a dedicated arm to directly identify demand effects, we incorporated three features to mitigate this risk. First, we introduced a four-month delay between the information treatment (June 2020) and our primary trust outcomes (October 2020). This temporal separation reduces the salience of the treatment when outcomes are measured. Second, we employed an obfuscated measurement strategy (Haaland et al., 2023). Trust in the CDC was measured four months after treatment, embedded within a battery of seven institutions presented in random order — the scientific community, banks, hospitals, health insurance companies, the police force, the private sector, and the CDC — making it difficult for respondents to identify it as the institution of interest. The follow-up knowledge question asked respondents to rank the U.S. relative to other countries in COVID-19 deaths per capita, a question that could not be answered using the CDC link, further obscuring the connection to the initial intervention. Third, we concluded the treatment wave with questions about the George Floyd protests and the Black Lives Matter movement — a politically salient topic unrelated to the pandemic — so that respondents’ final impression of the wave was not focused on COVID-19 content. Further, the intervening period coincided with the 2020

⁷Version 1.4 (Carril, 2017). Following Kerwin et al. (2024), we assess balance using an omnibus joint F-test; results are reported in Figure A6. We cannot reject balance across treatment arms ($p = 0.99$). The standard balance table is in Table A13 in the Appendix.

⁸Some studies advocate for incentive-compatible mechanisms to elicit accurate knowledge. We opted against financial rewards to maintain the voluntary nature of engagement. See Belot and Briscese (2023) on the importance of non-incentivized self-selection into contact.

⁹We did not pre-register the analysis plan for this experiment due to an oversight before the trial launch. For the second study, we pre-registered a detailed analysis plan and applied identical procedures, decision rules, and analytical methods to both studies, strengthening the reliability of findings from the first experiment.

Presidential election campaign, keeping respondents’ attention focused on politically polarizing events unrelated to our study.

An important consideration for external validity is the messenger’s role (Alatas et al., 2024). Information in our experiment was delivered through a survey administered by NORC, a highly reputable and nonpartisan research organization. Because respondents likely perceive NORC as a neutral and credible source, treatment effects may exceed those that would arise from direct institutional communication. While this limits what we can infer about institutions’ own capacity to rebuild trust through direct messaging, it mirrors many real-world contexts in which institutional information reaches citizens through trusted intermediaries — news outlets, public health campaigns, or community organizations. In our second experiment, we further address messenger and demand concerns by varying informational content while holding the institutional source constant, and continue to find significant effects.

Finally, although our meta-cognition measure — asking respondents to reflect on how their trust has changed since before the pandemic — could itself heighten demand effects by making belief change salient, we find a strong correlation between self-reported and observed trust changes across institutions, suggesting that reported increases reflect genuine shifts rather than reporting artifacts.

Outcomes and empirical model. We track four outcomes: (i) engagement with the treatment; (ii) response accuracy; (iii) trust in public institutions; and (iv) behavioral change. In the treatment wave, we tracked only the first two and deliberately avoided trust questions to prevent anchoring or priming. In wave 7 (October 2020), we measured institutional trust through perceived politicization,¹⁰ which we view as a primary determinant of institutional trust: when citizens perceive an institution as highly politicized, they question its objectivity and ability to serve the public interest without bias (Braithwaite and Levi, 1998). Perceived politicization was measured across the same seven institutions as above, presented in random order. After a set of unrelated questions to minimize consistency effects, respondents were asked to reflect on changes in their trust beliefs since before the pandemic.¹¹ This meta-cognitive measure captures respondents’ change in trust, along with awareness of their own belief updating. For behavioral outcomes, we elicited self-reported compliance with CDC-recommended practices, focusing on past behaviors rather than intentions to reduce social desirability bias. We asked whether respondents had received a COVID-19 test, whether they had tested positive (to assess whether increased testing reflected higher virus exposure rather than compliance), and which preventive actions they regularly adopted.¹²

¹⁰*Currently, how independent versus politicized do you think the following institutions are?* Response options: (1) Not influenced by politics at all; (2) Not too influenced by politics; (3) Somewhat influenced by politics; (4) Very influenced by politics.

¹¹*Thinking about your confidence in the institutions before the COVID-19 crisis started and today, to what extent would you say your views have changed?* Response options: (1) I have a worse opinion than before; (2) My opinion did not change; (3) I have a better opinion than before. The institution list included the same seven institutions, plus the President, and the U.S. Senate and Congress.

¹²Options: (1) I do not leave home unless strictly needed; (2) I keep a safe distance from others; (3) I try to avoid touching objects when outside; (4) I wash my hands thoroughly and frequently; (5) I cover my mouth in public with a mask or scarf; (6) None of the above.

Mediators and covariates. Our key mediator is media diet, which we argue shapes institutional trust by reinforcing the ideological cues that behaviorally inattentive individuals rely on when assessing unfamiliar institutions. We constructed a media slant index following the methodology described in Appendix A4.0.2. Respondents who consume slanted media exhibit different pandemic perceptions, lower knowledge accuracy, and distinct trajectories of institutional trust over the course of the crisis. We also collected four measures of pandemic exposure that serve as additional mediators and controls: (i) direct impact, either financial (loss of more than 20% of income between February and October 2020) or health-related (hospitalization of the respondent or someone close to them with COVID-19); (ii) indirect impact, economic (weekly variation in consumer expenditures in the respondent’s county) or health-related (log of the COVID-19 case counts in the respondent’s county). Construction details are in Appendix A4.0.2. We estimate a standard OLS regression with survey weights to maintain representativeness despite attrition:

$$\text{outcome}_{izw} = \alpha + \beta \text{ Treated}_i + \gamma \text{ Covid19}_{izw} + \theta \text{ News}_i + \psi \text{ Demographics}_i + \epsilon_{izw} \quad (3)$$

where outcomes are measured in wave w for individual i living in zip code z . Standard errors are robust to heteroskedasticity and clustered at the primary sampling unit to account for the survey design.

3.4 Results

3.4.1 Knowledge Accuracy and Trust

Engagement with the knowledge question was significantly higher in the treatment group than the control group — a 20% increase from a baseline of just under one minute, as tracked by a hidden embedded timer (Table A14 in the Appendix), providing support for H1.¹³ This additional time likely reflects respondents clicking on the CDC link, consistent with equation (1): reducing search costs c_i increased attention m^* to a high-variance piece of information — the COVID-19 death rate, arguably the most salient indicator of the pandemic’s severity — that could inform respondents’ decisions. We use a stringent accuracy measure, flagging as correct only responses where both the national and state death estimates matched CDC figures. In line with our second hypothesis, Table 1 shows that the treatment increased the probability of a correct answer by 10 percentage points (38% relative to a baseline of 28%). In the control group, 24% underestimated and 26% overestimated state deaths; in the treatment group, these shares fell to 16% and 19%, respectively. The pattern is similar for national deaths: underestimates fell from 23% to 19% and overestimates from 26% to 21%. The treatment also had lasting effects on knowledge: treated respondents were 10% more likely to correctly rank the U.S. in terms of COVID-19 deaths per capita four months later. Turning to trust, the treatment significantly increased the perception that the CDC was politically independent by 6 percentage points (control-group mean: 8.0%), and treated re-

¹³At the time of implementation, NORC did not have the capability to track click-through rates, which we do in our second experiment.

spondents were 6.2 percentage points more likely to report that their confidence in the CDC had increased since the start of the pandemic (control-group mean: 12.3%), supporting our third hypothesis. An instrumental variable approach confirms that these effects are driven by knowledge accuracy: respondents who correctly estimated COVID-19 deaths were substantially more likely to report higher trust in the CDC (Table A15 in the Appendix). The first stage is strong when conditioning on demographics ($F = 13.5$), though weaker without controls ($F = 6.24$). Because Table 1 tests four outcomes, Appendix A3 reports inference adjusted for multiple testing. The June knowledge effect survives every adjustment. With controls, the two trust effects have sharpened false-discovery-rate q-values of 0.04 within the family of four outcomes and Romano-Wolf family-wise p-values of 0.06; the long-run knowledge effect has a q-value of 0.05 but a Romano-Wolf p-value of 0.16, and we therefore treat it as suggestive (Table A6).

Table 1: Treatment effects on knowledge in the short- and long-run, and on trust beliefs.

VARIABLES	June (w6)		October (w7)					
	(1) Correct death est.	(2) Correct death est.	(3) Correct US death rank	(4) Correct US death rank	(5) Incr. trust in CDC	(6) Incr. trust in CDC	(7) CDC not infl. by politics	(8) CDC not infl. by politics
Informative T	0.101*** (0.0324)	0.125*** (0.0286)	0.0406 (0.0336)	0.0539* (0.0305)	0.0637* (0.0325)	0.0618** (0.0292)	0.0545* (0.0301)	0.0599** (0.0269)
Constant	0.290*** (0.0217)	0.334** (0.143)	0.528*** (0.0294)	1.039** (0.443)	0.123*** (0.0176)	0.391 (0.281)	0.0795*** (0.0194)	-0.0399 (0.188)
Observations	1,147	1,147	977	977	978	978	990	990
R-squared	0.022	0.161	0.002	0.082	0.014	0.088	0.021	0.136
Demographics	No	Yes	No	Yes	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.330	0.330	0.550	0.550	0.162	0.162	0.113	0.113

Notes: Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$. OLS coefficients. Cols (1–2): correct estimation of COVID-19 deaths in the U.S. and respondent’s state, June; cols (3–4): correct ranking of the U.S. by COVID-19 deaths per capita, October; cols (5–6): self-reported increase in confidence in the CDC since the start of the pandemic, October; cols (7–8): CDC perceived as not influenced by politics, October. All regressions use population survey weights. Control variables are described in the Appendix.

Meta-cognition. A potential concern with survey experiments is that self-reported changes in trust may not reflect actual belief updating. To assess the validity of this meta-cognitive measure, we examine whether self-reported changes predict observed variations in trust across a set of institutions tracked between April and October. We find a strong and consistent relationship between the two (Tables A27 and A28 in the Appendix), indicating that respondents’ reports align closely with realized changes in trust. This is consistent with prior evidence that self-reported trust captures meaningful underlying beliefs (e.g., Sofianos (2022)), and extends it by showing that meta-cognitive measures are informative not only about levels, but also about changes. More generally, meta-cognitive measures can serve as a useful complement for tracking within-individual belief changes when panel data are limited.

Prior beliefs and who searches. News consumption proxies the stock of beliefs a respondent brings to the question, and trust in the scientific community proxies the perceived cost of engaging with the institution. Both predict knowledge accuracy strongly. Respondents who follow U.S. news for at least thirty minutes a day are 9.7 to 12.4 percentage points more

likely to answer correctly, against a 39 percent base rate, while those reporting at most “some confidence” in the scientific community in wave 1 — two months before the experiment, and so pre-treatment — are 7.2 to 12.2 points less likely to do so (Table A2 in the Appendix). The gap in the time respondents spend searching appears only where searching was costly. In the control group, which faced the higher search cost, low-trust respondents spent 13.9 fewer seconds on the question; among the treated, who were offered the link, the gap disappears, and the interaction is 31.8 seconds with a standard error of 13.7 ($p = 0.022$). This is what the model predicts: when the search cost is high it is borne by those for whom it is perceived as low, and lowering it brings the others in. Appendix A2 reports the full results, including the parallel evidence from Study II, and discusses their limits — news consumption and trust are chosen rather than assigned, and the two are sufficiently correlated that they cannot be cleanly separated.

3.4.2 Behaviors

We hypothesized that rebuilding institutional trust would increase compliance with public health guidelines — an expectation supported by the literature showing that citizens who perceive the government as trustworthy are more likely to comply with its recommendations (Tyler, 1990; Braithwaite and Levi, 1998). The evidence is mixed (Table A16 in the Appendix). Treated respondents were 20% more likely to have taken a COVID-19 nasal swab test (6.6 percentage points, from 28.1% in the control group). This effect was not driven by a higher rate of positive tests (Table A18 in the Appendix), suggesting the increase reflects greater compliance with public health recommendations rather than illness. Adjusted for the four behavioural outcomes tested together, the testing effect has a sharpened q-value of 0.06 but a Romano-Wolf p-value of 0.20 (Table A6 in the Appendix), which is why we describe the behavioural evidence as suggestive. Whether driven by improved knowledge of the pandemic’s severity or by higher trust in the CDC, respondents acted in line with institutional guidance. For other preventive behaviors, the treatment increased the probability of adopting at least one recommended behavior, though this effect is concentrated among a small number of individuals and should be interpreted cautiously; we find no effect at the intensive margin. By October 2020, most pandemic restrictions had been lifted, and masks were no longer mandatory, which likely diminished the relevance of these measures and explains the absence of detectable effects. Of the behaviors still actively recommended at that time, COVID-19 testing was the most consequential for limiting viral spread, making it the most meaningful behavioral outcome to shift. An IV regression analogous to the one above suggests effects are driven by knowledge accuracy, though TSLS coefficients are significant only with demographic controls (Table A17 in the Appendix).

3.4.3 Robustness Checks and Alternative Mechanisms

We summarize here the additional analyses reported in full in the Appendix.

Placebo outcomes. We measured trust across seven institutions that were shown in random order to survey respondents and found no significant treatment effects on any institution other than the CDC (Table A19 in the Appendix), confirming that the trust improvement

was targeted rather than a broad attitudinal shift. The only exception is a modest spillover to the scientific community, which is unsurprising given the CDC’s identity as a scientific institution and suggests that trust gains can extend to closely related credible institutions.

Attrition. Approximately 75% of respondents completed the final survey wave, with an attrition rate of 17% between the last two waves. All analyses use survey weights that adjust for attrition to preserve population representativeness. Nonetheless, attrition between the experimental and outcome waves could bias estimates if it differs systematically by treatment status. We rule out this risk in three ways. First, attrition rates do not differ significantly across treatment arms, either unconditionally ($p = 0.84$) or conditional on baseline covariates ($p = 0.94$). Second, although baseline covariates jointly predict attrition (F-test = 5.56, $p < 0.001$), consistent with selective panel retention in survey data, we find no evidence that treatment assignment alters these patterns (joint F-test of treatment $\times$ covariate interactions = 1.02, $p = 0.42$). Third, treatment effects generally do not vary with characteristics that predict attrition (Tables A20-A23 in the Appendix). The only exception is the perceived political independence of the CDC, where estimated effects are larger among Hispanic respondents ($N=133$; treatment effect: 0.15, $p < 0.05$) than among non-Hispanic respondents (0.034, $p = 0.2$). This pattern does not extend to other trust or behavioral outcomes, and the Hispanic subsample is relatively small. Notably, treatment reduced attrition among Hispanic respondents by 9 percentage points, suggesting the intervention may have been particularly engaging for this group. Together, weighting, balance tests, and the absence of treatment–covariate interactions in attrition suggest that any selective panel retention is orthogonal to treatment assignment, reducing concerns that differential selection into the outcome sample biases the estimated treatment effects.

Interpersonal trust. Controlling for baseline interpersonal trust reported in wave 1 leaves the treatment effect unchanged (Table A24 in the Appendix), ruling out the possibility that our results are driven by generalized social trust.

Perceived politicization and trust In line with the literature (Bolsen and Druckman, 2018; McLaughlin et al., 2021; Van Scoy et al., 2021; Clark et al., 2023), trust erosion tracked closely with perceived politicization. Respondents who believed an institution was affected by politics reported lower confidence in it relative to the first wave and were more likely to report having decreased their trust over time (Table A11 in the Appendix).

Heterogeneity. We find no robust heterogeneity in treatment effects on trust by media consumption, political affiliation, education, COVID-19 exposure, or baseline trust in science. The one interaction on a trust outcome that approaches conventional significance is that between treatment and Republican-leaning news consumption on perceived political independence, which is negative (-0.099 , s.e. 0.060, $p = 0.10$; Table A29 in the Appendix), so that the point estimate of the effect for this group is close to zero. The interactions for Republican- and Democratic-leaning consumers cannot be distinguished from each other ($p = 0.24$), and no comparable pattern appears for self-reported confidence, so we read it as suggestive at most. Significant heterogeneity emerges only for long-term knowledge accuracy: the treatment effect is 26 percentage points larger ($p < 0.001$) among Republican-leaning

news consumers than Democratic-leaning ones and 11 percentage points larger ($p < 0.1$) among respondents without a college degree, both groups exhibiting substantially lower baseline knowledge in the control condition. The absence of robust partisan heterogeneity in trust is notable given evidence on motivated reasoning, which suggests that individuals consuming politically opposed media may resist institution-sourced information.

Several features of our design may help explain this pattern. The intervention delivered factual, verifiable metrics rather than normative policy recommendations and was administered through NORC, a nonpartisan survey organization (messenger effect), potentially reducing resistance at the point of intake. In addition, the near-zero search costs embedded in the intervention may have limited opportunities for motivated avoidance. While we cannot distinguish empirically among these mechanisms, the results highlight the importance of informational design in shaping the effectiveness of institutional communication.

Complementary evidence from the control group is consistent with the search-cost mechanism emphasized in our model. Among untreated respondents, those consuming Republican-leaning media — who report lower baseline trust in the scientific community — spent significantly less time on the knowledge question than those consuming Democratic-leaning media ($p = 0.08$ without controls and $p = 0.02$ with them; Table A30 in the Appendix). This pattern suggests less voluntary engagement with institutional information among lower-trust individuals. The corresponding association with knowledge accuracy is directionally similar but not robust to covariate adjustment, likely reflecting confounding by socioeconomic characteristics. Taken together, these results are consistent with a framework in which trust facilitates voluntary information acquisition, whereas the experimental reduction in search costs attenuates reliance on such trust-driven engagement. Whether repeated exposure to independent information sources may also affect media consumption over time remains an open question.

Alternative specifications. Results are robust to a probit model (Table A25).

4 Study II: The Importance of Information Value

Our first study demonstrated that reducing search costs for institution-sourced information encourages citizens to engage with it, update their factual knowledge, and increase trust, with suggestive behavioral improvements. A second study provides additional evidence on the mechanism: does increased trust stem from mere exposure to institutional communication, or does the type of trustworthiness signal matter?

4.1 Data

We partnered with NORC to administer a survey experiment to the GSS AmeriSpeak panel in August 2024, four years after our original study, leveraging the Time-sharing Experiments for the Social Sciences (TESS) program.¹⁴ We recruited 1,630 respondents and implemented

¹⁴TESS is an NSF-funded initiative that fields survey experiments on NORC’s AmeriSpeak Panel at no cost to researchers, though it limits investigators to six items per study.

a three-arm design that replicates the key features of the first study while introducing variation in information value to isolate the mechanism. We intentionally timed this experiment ahead of the 2024 Presidential elections to mirror the politically polarized context of our first study, and to examine whether the same intervention produces similar effects outside a public health crisis and across a different institution. We chose to focus on trust in the U.S. Environmental Protection Agency (EPA). EPA policies directly affect Americans’ daily lives in ways that are personally salient — unlike institutions such as NASA or NOAA, whose mandates are less immediately felt — and compliance with EPA-recommended behaviors (e.g., reducing one’s carbon footprint) generates positive welfare spillovers comparable to following CDC guidance during the pandemic. Like the CDC in 2020, the EPA was subject to significant politically polarizing attacks at the time of our study, with Republican leadership calling for the resignation of the EPA’s administrator and seeking a Supreme Court injunction to block new environmental protection rules, making it a suitable institutional context for our complementary study. The survey comprised six questions. The first assessed respondents’ baseline familiarity with the EPA,¹⁵ allowing us to test whether, consistent with our behavioral inattention hypothesis, unfamiliarity correlates with lower baseline trust. The second introduced our randomized intervention (detailed below). Questions three and four measured two dimensions of trust: perceived benevolence¹⁶ and perceived politicization,¹⁷ both elicited on 7-point Likert scales following Stantcheva (2023), from 1 (“Always worsen” / “Completely politicized”) to 7 (“Always improve” / “Completely independent”), with 4 as the neutral midpoint. The same questions were asked about the IRS as a placebo. Question five elicited respondents’ support for EPA involvement in a hypothetical community research scenario.¹⁸ Since government agencies cannot receive donations, this willingness-to-recommend question serves as our behavioral proxy for institutional support. The sixth and final question asked about preferred news sources, enabling us to construct a media slant index using the same methodology as the first study.¹⁹

¹⁵“Look at the list of government agencies below and select the ones that you are familiar with — that is, the ones you know something about their mission, role, and responsibilities.” Respondents chose from a randomized list including: EPA, CDC, NOAA, National Park Service, FDA, IRS, Federal Laboratory Consortium for Technology Transfer, and Open World Leadership Center.

¹⁶“Do you think the people running the following institutions recommend policies and regulations that improve or worsen your life and the lives of your loved ones?”

¹⁷“How politicized versus independent do you think the following institutions are?”

¹⁸“Imagine your state is planning a study to measure air pollution levels in our community. They are asking residents like you which organization they trust most to lead this important project. What is the percent chance on a scale from 0% (absolutely no chance) to 100% (absolutely certain) that you would recommend the Environmental Protection Agency (EPA)?” We use Wiswall and Zafar (2015)’s approach to aid interpretation of percentage responses, providing illustrative anchors: e.g., 2–5% as “almost no chance,” 47–52% as “pretty even chance,” 95–98% as “almost certain.”

¹⁹After the last question, all respondents were shown links to the EPA’s Facebook, X and Instagram pages and invited to follow the agency; whether they clicked any of them was a pre-registered secondary outcome, intended as a low-cost measure of willingness to keep receiving the institution’s information. We find nothing on it: 9.3 percent of respondents clicked a link, the share is not higher in T1 than in T3 (1.5 percentage points, s.e. 1.8), and among the two link arms it does not differ between respondents who did and did not open the treatment link (−1.0 points, s.e. 2.0). The one difference across arms is that T2, which received no link, clicked 4.3 points more often than T3 ($p = 0.01$), a contrast the design gives us no reason to expect and which is not distinguishable from T1 ($p = 0.12$); we do not interpret it. Table A36 in the Appendix reports the estimates.

4.2 Experimental Methodology

Regardless of prior EPA familiarity, all respondents first received a brief description of the agency.²⁰ This provided a common baseline understanding of the EPA’s mission for all participants. We then tested factual knowledge using a three-arm randomization. In both treatment arms (T1 and T2), respondents were shown a two-sentence description of the Air Quality Index (AQI) — the EPA’s real-time measure of air pollution — and asked to report the current AQI in their zip code using a slider from 0 to 500.²¹ In T1 (low search cost, high-value information), respondents were additionally offered a direct link to the EPA’s air quality tracker, mirroring the design of the first experiment. The control group (T2) received the AQI question without a link, facing higher search costs. In T3 (low search cost, low-value information), respondents were instead asked about the location of the EPA regional office responsible for their state, with an equivalent link provided.²² This design allows us to vary two parameters of equation (1) independently: the search cost c_i (T1 vs. T2) and the informational value x_i (T1 vs. T3), while holding the institutional source constant across all arms. Crucially, because all respondents are exposed to information about the same institution, differences in outcomes across arms cannot be attributed to messenger endorsement or general experimenter demand effects — they must reflect the content and perceived usefulness of the information itself. As in the first experiment, respondents were not incentivized to consult the EPA website and could freely choose whether to do so. Respondents could also leave the factual question blank, and 14 percent did: 3 percent in T2, against 20 percent in T1 and 19 percent in T3, where non-response is concentrated among those who opened the link (28 and 34 percent of clickers, against 6 and 2 percent of non-clickers), most plausibly because the link opened in a new window and some respondents did not return to the question. This is why the knowledge columns of Table 2 have 1,400 observations and the trust columns about 1,600; because leaving the question blank is a post-treatment choice, the trust outcomes are estimated on the full sample rather than conditioned on answering, and the T1–T3 trust effects are, if anything, slightly larger among those who answered. Our measure of information consumption — the attention m_i^* of the model — is whether the respondent answered the factual question correctly. We record whether respondents opened the link, but a click is a weak measure of consumption: we cannot observe what happens after the page opens, and a respondent may close it immediately. Answering correctly, by contrast, requires having actually taken in the number, whether from the EPA’s page or elsewhere, so we use it as the primary measure throughout and treat clicking as a supplementary one. We verified response accuracy by merging respondents’ zip codes and survey completion times with publicly available EPA AirNow data, and cross-referenced T3 responses with EPA regional office assignments by state. Because AirNow displays the AQI of whichever pollutant is highest rather than PM2.5 specifically, we score respondents against the headline AQI

²⁰“The U.S. Environmental Protection Agency (EPA) is the federal government agency responsible for implementing policies that protect human health and the environment, such as reducing emissions.”

²¹“EPA also manages the NowCast Air Quality Index (AQI), which tracks levels of air pollution in real time across the country. The higher the NowCast AQI score, the worse the air quality and the greater the health concern. To the best of your knowledge, what is the current level of pollution in your ZIP code?”

²²“EPA has multiple regional offices to address health and environmental concerns in specific areas of the country, including where you live. To the best of your knowledge, in which city is the EPA regional office that is responsible for your state of residence?” (see Figure A8)

for their reporting area, and we convert survey timestamps to GMT before matching. Our pre-analysis plan scored an answer as correct if it fell within ± 10 of the benchmark; in the main text we use the tighter ± 5 window, for a reason that matters to the interpretation. T3 is scored on a question with no bandwidth, so its accuracy is fixed regardless of the window, whereas T1 and T2 accuracy rises mechanically as the window widens. At ± 5 , T1 and T3 answer their questions correctly at the same rate, which lets us compare trust across the two arms holding accuracy fixed; at ± 10 , T1 exceeds T3. We report the pre-registered window in the Appendix (Table A33) and note that neither the search-cost result nor any trust outcome depends on the choice.²³ We estimate a standard OLS regression for individual i in zip code z , excluding T3 so that the T1 coefficient captures the effect of low-cost, high-value information relative to low-cost, low-value information:

$$\text{outcome}_{iz} = \alpha + \beta_1 T1_i + \beta_2 T2_i + \theta \text{News}_i + \psi \text{Demographics}_i + \epsilon_{iz} \quad (4)$$

Standard errors in Table 2 are classical; the Appendix tables state where heteroskedasticity-robust standard errors are used instead.

4.3 Results

We examine two dimensions of trust — perceived benevolence and political independence — and assess whether search costs and information value shape them differently. Consistent with H2, both low-search-cost conditions (T1 and T3) yield significantly higher knowledge accuracy than the no-link condition (T2), which reduces accuracy by approximately 30 percentage points relative to T3 (64 percent, cols. 1–2 of Table 2) and by 29 points relative to T1 (p-val T1 = T2, Table 2). The two link arms answer their respective questions correctly at the same rate, 45.9 percent in T1 and 47.9 percent in T3, so in what follows the T1–T3 comparison is one across arms with equal accuracy and different content. Respondents in T1 are also 12 percentage points more likely to click on the provided link than those in T3 ($p < 0.001$; Table A35 in the Appendix),²⁴ so more informative cues attract greater attention even when the cost of reaching them is identical.

Turning to trust, we follow the pre-analysis plan and binarise each outcome: agreeing

²³AirNow’s hourly monitor files report NowCast values: recomputing the AirNow PM NowCast from the reported hourly concentrations reproduces the published PM2.5 index in 97.9 percent of 318,242 monitor-hours. For each reporting area and hour we take the maximum across monitors of each pollutant’s index and then the maximum across pollutants, which is the pollutant AirNow flags as primary in every reporting area we can check against its own reporting-area files. In our field window PM2.5 is the headline pollutant in only 62 percent of area-hours, with ozone accounting for a further 33 percent and rising to about half in the late afternoon. Survey timestamps are recorded in a single fixed United States zone, which the survey documentation does not identify; we select the offset empirically (UTC−4), which maximises agreement with the answers of respondents known to have opened the AirNow page. Among those respondents the benchmark reproduces the reported value exactly in 18.4 percent of cases and to within ± 10 in 74.8 percent, against 6.0 and 58.0 percent under a PM2.5-only match using unconverted timestamps. Because AirNow publishes reporting-area values through a separate agency-reported process, a reconstruction from the monitor files is close to but not identical with the displayed number, so the exact-match rate is a lower bound on the share of respondents who consulted the page.

²⁴We do not compare time spent across treatment arms, as respondents completed different tasks.

that the EPA improves lives and that it is politically independent (above the midpoint of the 7-point scale), and reporting a likelihood above 50 of recommending the EPA to lead a community study. Providing access to local AQI data raises the probability that respondents view the EPA as politically independent by 6.2 percentage points relative to the office-location link ($p = 0.013$), a 34 percent increase on the T3 base of 18.0 percent, and by 4.6 points with the full control set ($p = 0.063$; cols. 5–6). It also raises the share who agree that the EPA improves lives by 5.0 points (cols. 3–4) and the share who would back the EPA to lead a study by 6.3 points (cols. 7–8); both are significant without controls and attenuate with them. Because three trust outcomes are tested at once, Appendix A3 reports adjusted inference: without controls the independence effect has a sharpened false-discovery-rate q -value of 0.04 and a Romano-Wolf family-wise p -value of 0.045, and the support effect 0.04 and 0.07; with the full control set no adjusted value falls below 0.17 (Table A7). The same outcomes on their original scales give the same pattern (Table A32 in the Appendix). The T1–T3 contrast is the central test of H3 in our design, and it is worth being explicit about why. Both arms are given a link to the same agency, in the same format, at the same point in the survey, and both therefore face the same search cost c_i ; they differ only in whether the information on the other end of the link bears on a decision the respondent might actually take. The differences just reported are therefore attributable to the value of the information rather than to the cost of reaching it, to the identity of the messenger, or to a general willingness to comply with the survey, and they do not rest on any auxiliary split of the sample. A plausible interpretation is that real-time air quality data signals greater technical rigor and accountability in EPA operations, strengthening confidence in the agency. This is the distinction between relevance and distortion drawn at the end of Section 2: consulting a verifiable number is informative about whether the agency’s output is politically distorted, which is what the independence question elicits, and less so about whether its work is relevant, which is closer to what the benevolence question elicits.

Table 2: Results on trust outcomes.

	(1) Knowledge accuracy (± 5)	(2) Knowledge accuracy (± 5)	(3) EPA improves lives (0/1)	(4) EPA improves lives (0/1)	(5) EPA indepen- dent (0/1)	(6) EPA indepen- dent (0/1)	(7) Support for EPA (0/1)	(8) Support for EPA (0/1)
T1: Low search costs & high value info	-0.0187 (0.0311)	-0.0276 (0.0312)	0.0498* (0.0296)	0.0233 (0.0286)	0.0615** (0.0248)	0.0456* (0.0245)	0.0634** (0.0300)	0.0478 (0.0291)
T2: High search costs & high value info	-0.305*** (0.0292)	-0.300*** (0.0293)	0.0177 (0.0292)	0.0117 (0.0280)	0.0335 (0.0244)	0.0252 (0.0240)	0.0393 (0.0296)	0.0331 (0.0287)
Constant	0.487*** (0.0216)	0.465*** (0.0935)	0.596*** (0.0207)	0.394*** (0.0871)	0.180*** (0.0174)	0.237*** (0.0749)	0.591*** (0.0210)	0.607*** (0.0895)
Observations	1,400	1,400	1,626	1,626	1,613	1,613	1,580	1,580
R-squared	0.097	0.127	0.015	0.117	0.006	0.067	0.004	0.094
Demographics	No	Yes	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes	No	Yes
p-val T1 = T2	0.000	0.000	0.280	0.686	0.260	0.406	0.426	0.616
Mean dep. var.	0.356	0.356	0.609	0.609	0.207	0.207	0.622	0.622

Notes: Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$. OLS coefficients. Excluded group: low search cost, low-value information (T3). Cols (1–2): knowledge accuracy, scored to within ± 5 of the headline AQI (T1, T2) or as the correct regional-office city (T3); the pre-registered ± 10 window is in Table A33 in the Appendix; cols (3–8): trust outcomes binarised as specified in the pre-analysis plan — agrees the EPA improves lives (cols 3–4) and that it is politically independent (cols 5–6), both above the midpoint of the 7-point scale, and reports a likelihood above 50 of recommending the EPA to lead a community air quality study (cols 7–8). The same outcomes on their original scales are in Table A32 in the Appendix. Standard errors and the p-val row are classical throughout. Controls are described in the Appendix.

The p-values in the second-to-last row show that T1 and T2 never differ significantly

in trust outcomes, despite T1 producing substantially higher knowledge accuracy.²⁵ Two complementary explanations account for this pattern. First, in contrast to the CDC experiment — where searching for COVID-19 mortality data could yield heterogeneous sources — respondents in T2 were explicitly instructed to find the EPA’s AQI score, making it likely that any search led directly to the same EPA page as T1. The search cost difference thus compressed differences in information exposure between arms. Second, the AQI question itself — by explicitly referencing EPA’s real-time pollution monitoring — may have signaled institutional competence and relevance to all high-value information respondents, regardless of whether the link was provided. This interpretation is supported by the T3 results: when the question referenced lower-value information (office locations), trust effects were smaller or null even with a link, consistent with the information content itself — not merely the link — driving the trust response.

4.3.1 Mechanisms: Engagement, Accuracy, and Trust

A natural concern with the T1–T3 comparison is that respondents in T1 might trust the EPA more simply because more of them answered correctly, so that the trust effect reflects knowledge rather than the value of the information consulted. At the ± 5 window used in the main text this cannot be the explanation: T1 and T3 answer their respective questions correctly at indistinguishable rates (45.9 and 47.9 percent; Table 2, cols. 1–2). The two arms differ in what they were shown, not in how many got it right.

What respondents see matters. Table 3 compares respondents who opened the link in T1 with respondents who opened the link in T3 — the same action, taken at the same point in the survey, leading to different pages. Among those who clicked, respondents shown the air quality index are 5.9 percentage points more likely to view the EPA as politically independent than respondents shown the regional office (col. 4, $p = 0.058$; 7.0 points without controls, $p = 0.026$), and 0.26 points higher on the seven-point scale ($p = 0.010$; Table A34 in the Appendix). Among those who did not click, the T1–T3 difference is smaller and not significant. Opening the EPA’s page raised perceived independence when the page carried something useful and not when it carried an address. The two groups’ differences cannot themselves be distinguished statistically, so we do not claim the effect is confined to clickers; but holding the act of consulting fixed and varying only the content, it is the informative content that goes with higher trust.

Two qualifications apply. Opening the link is a post-treatment choice, so every coefficient in Table 3 is a descriptive association within arms; the causal evidence remains the intention-to-treat comparison in Table 2. And clicking is a coarse measure of consumption — it records that the page was opened, not whether anything was read, and a T1 respondent who did not click may still have consulted the index on a phone — which is why answering correctly, not clicking, is our primary measure of consumption in Table 2, and why the click-based contrasts here should be read as attenuated.

²⁵While our pre-analysis plan specified comparisons against the no-link control, the similarity of trust effects for T1 and T2 led us to focus on the theoretically more informative T1 vs. T3 comparison, which was also pre-specified.

Table 3: Trust and opening the link, by treatment arm.

VARIABLES	(1) EPA improves lives	(2)	(3) EPA independent (binary)	(4)	(5) Support for EPA	(6)
T1: Low search costs & high value info	0.0578 (0.0485)	0.0382 (0.0475)	0.0532 (0.0410)	0.0422 (0.0418)	0.0387 (0.0487)	0.0241 (0.0483)
Clicked link	0.0788* (0.0420)	0.0700* (0.0407)	-0.0216 (0.0329)	-0.0158 (0.0333)	-0.0008 (0.0426)	-0.0079 (0.0415)
Clicked $\times$ T1	-0.0263 (0.0615)	-0.0305 (0.0599)	0.0170 (0.0517)	0.0165 (0.0519)	0.0374 (0.0621)	0.0390 (0.0610)
Constant	0.552*** (0.0316)	0.245** (0.105)	0.191*** (0.0252)	0.269*** (0.0874)	0.592*** (0.0317)	0.503*** (0.115)
T1 vs T3, among those who clicked	0.0315 (0.0379)	0.0076 (0.0369)	0.0702** (0.0314)	0.0587* (0.0309)	0.0761** (0.0386)	0.0631* (0.0374)
T1 vs T3, among those who did not	0.0578 (0.0485)	0.0382 (0.0475)	0.0532 (0.0410)	0.0422 (0.0418)	0.0387 (0.0487)	0.0241 (0.0483)
Observations	1,075	1,075	1,067	1,067	1,049	1,049
R-squared	0.020	0.123	0.009	0.070	0.006	0.093
Demographics	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.611	0.611	0.206	0.206	0.619	0.619

Notes: Standard errors in parentheses, robust to heteroskedasticity. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

OLS on the two arms that were offered a link (T1 and T3), the only arms in which clicking is recorded. Excluded arm: T3. The rows below the line report the T1–T3 difference among respondents who opened the link and among those who did not; the latter equals the coefficient on T1. Outcomes are binary, as in Study I: agrees that the EPA improves lives and that it is politically independent (above the scale midpoint), and would support the EPA leading a community air quality study (> 50); original scales in Table A34. Clicking is a post-treatment choice, so these are descriptive associations. Note also that the air quality index is available from many weather applications, so a T1 respondent who did not open the EPA link may still have consulted the number elsewhere. Produced by TESS/Codes/EPA/11.mechanism_tables_w5.do.

4.3.2 Robustness Checks

Knowing the EPA. Approximately 30% of respondents did not identify the EPA as a familiar institution. As shown in Table A37, familiarity with the EPA is positively correlated with perceived benevolence and with support for its involvement in a community study, but negatively correlated with perceived political independence: respondents who know the EPA rate it 0.15 to 0.18 points more politicized on the seven-point scale. The two patterns are not in tension once one asks who the familiar respondents are. Familiarity is not the same as having consumed useful information from the agency. For a partisan respondent, knowing the EPA can mean knowing it as an object of political controversy, and in August 2024 the agency was under sustained attack from Republican leadership. Table A38 splits the association by news slant. Among consumers of Republican-leaning news, knowing the EPA is associated with a 0.64 point lower independence rating (s.e. 0.17), against 0.22 points (s.e. 0.09) among consumers of balanced news; among consumers of Democratic-leaning news the sign reverses, and familiarity goes with a 0.53 point *higher* rating (s.e. 0.19). In raw means, Republican-leaning respondents who know the EPA rate its independence at 2.5 on the seven-point scale, against 3.1 for Republican-leaning respondents who do not know it and about 3.5 for everyone else. The negative unconditional association therefore reflects the presence, among those who know the agency, of Republican-leaning respondents who

know it and distrust it, rather than a general tendency for familiarity to breed a perception of politicization. Consistent with the first study, Republican-leaning media consumers were also less likely to report knowing the EPA ($p < 0.10$; Table A39). Among respondents in the control group T2 (no link provided), knowing the EPA remains a strong predictor of trust even after controlling for knowledge accuracy, consistent with our theoretical prediction that familiarity lowers perceived information costs and reinforces trust independently of any single informational interaction.

Placebo outcomes. Treatment effects on IRS trust are not significant in any specification but one, and that effect disappears with controls (Table A40 in the Appendix), suggesting at most limited demand effects.

Alternative specifications. The trust outcomes on their original Likert and 0–100 scales (Table A32 in the Appendix) give the same pattern as the binarised versions in Table 2; a probit first stage and ordered probit for the Likert-scale outcomes are also consistent (Table A41).

EPA office location. The null trust effects in T3 are not driven by respondents discovering that the EPA has no office in their state and updating negatively. There is no evidence of this mechanism in the data (Table A42 in the Appendix).

Heterogeneity. We find no significant heterogeneity by news consumption, political affiliation, education, or age. Respondents in more polluted areas (zip codes with a headline AQI above 50) are somewhat more likely to answer correctly when treated, but the interaction is significant only once controls are included and is far smaller than the main treatment effect, and we find no corresponding heterogeneity in any trust outcome (Table A43 in the Appendix). We would in any case caution against reading this split as a test of the σ_i^2 channel. An AQI above 50 records how *severe* air quality is at the moment of the survey, whereas σ_i^2 is the variability of the index over time, which is what would make it worth consulting repeatedly; the two are close to orthogonal across the zip codes in our sample. Measuring that variability credibly requires several months of AirNow history rather than the field window itself, and we leave it to future work. Our evidence that information value matters therefore rests on the T1–T3 contrast, which is the cleaner test in any case: it holds the institution, the search cost, and the format of the prompt fixed, and varies only the value of the information. We cannot test whether respondents in the control group spent more time on the treatment question by news consumption, as we did not measure this dimension. We have access to the time spent on the survey overall, but this measure is noisier.

5 Conclusions

Trust in public institutions, essential for promoting behaviors that address societal challenges, has declined across Western democracies. This paper provides evidence that limited engagement with institution-sourced information contributes to this decline. Across two

theory-driven survey experiments conducted on nationally representative samples of Americans, we show that reducing the costs of accessing relevant information increases knowledge and raises institutional trust. When individuals engage with institution-provided content, they become more informed and report higher trust, with suggestive evidence of increased compliance with recommended behaviors. These effects persist over time and are stronger when the information is valuable. Taken together, our findings are consistent with a mechanism in which frictions in information acquisition limit exposure to credible institutional signals, reinforcing reliance on prior beliefs and heuristics. By lowering these frictions, even light-touch scalable interventions can shift both knowledge and trust, suggesting that part of the erosion in institutional confidence may reflect inattention rather than deeply entrenched preferences.

Importantly, both experiments were conducted in politically salient environments, where baseline attention to institutions may already be elevated. While this could, in principle, attenuate the marginal impact of reducing search costs, the consistency of our results across different institutional contexts and time periods nonetheless suggests that the underlying mechanism is not confined to moments of crisis. At the same time, our design relies on a trusted intermediary, NORC, a widely respected nonpartisan research organization, which may enhance the credibility of the information and, in turn, amplify the observed effects. While this mirrors real-world channels through which institutional information often reaches citizens, it remains an open question as to whether the magnitude of the effects we detect may differ when the messaging comes directly from a public institution. A large-scale field experiment, conducted in partnership with an institution, can provide answers to such important questions.

Future work could also examine the scalability of these interventions across diverse institutional and political environments. Extending the framework internationally would help clarify how cultural and educational contexts shape the impact of interventions. Additional research may also explore whether trust gains spill over into other domains of social capital and how responses vary when institutional information is more complex or cognitively demanding, helping to disentangle the roles of cognitive capacity and behavioral inattention in shaping trust (Nunnari et al., 2024; Dal Bó et al., 2018).

To conclude, the results of our studies suggest that institutional trust is not an immutable attitude, and that light-touch scalable digital interventions may help rebuild trust. Although political dynamics, elite discourse, and partisan news outlets will remain central, our results indicate that institutional outreach efforts can also help counteract the decline in public trust. The absence of robust partisan heterogeneity in treatment effects in both of our studies offers a cautiously optimistic insight that this type of intervention may yield benefits across the political spectrum.

References

Acemoglu, D., A. Cheema, A. I. Khwaja, and J. A. Robinson (2020). Trust in state and nonstate actors: Evidence from dispute resolution in pakistan. *Journal of Political Econ-*

omy 128(8), 3090–3147.

- Alabrese, E., F. Capozza, and P. Garg (2024). Politicized scientists: Credibility cost of political expression on twitter. *Working Paper*.
- Alatas, V., A. G. Chandrasekhar, M. Mobius, B. A. Olken, and C. Paladines (2024). Do celebrity endorsements matter? a twitter experiment promoting vaccination in indonesia. *The Economic Journal* 134(659), 913–933.
- Alesina, A. and G. Tabellini (2007). Bureaucrats or politicians? part i: A single policy task. *American Economic Review* 97(1), 169–179.
- Algan, Y., D. Cohen, E. Davoine, M. Foucault, and S. Stantcheva (2021). Trust in scientists in times of pandemic: Panel evidence from 12 countries. *Proceedings of the National Academy of Sciences* 118(40), e2108576118.
- Alsan, M. and M. Wanamaker (2018). Tuskegee and the health of black men. *The quarterly journal of economics* 133(1), 407–455.
- Anderson, M. L. (2008). Multiple inference and gender differences in the effects of early intervention: A reevaluation of the abecedarian, perry preschool, and early training projects. *Journal of the American statistical Association* 103(484), 1481–1495.
- Banerjee, A., M. Alsan, E. Breza, A. G. Chandrasekhar, A. Chowdury, E. Duflo, P. Goldsmith-Pinkham, and B. A. Olken (2024). Can a trusted messenger change behavior when information is plentiful? evidence from the first months of the covid-19 pandemic in west bengal. *Review of Economics and Statistics*, 1–33.
- Belot, M. and G. Briscese (2023). *Bridging America’s Divide on Abortion, Guns and Immigration: An Experimental Study*.
- Bénabou, R. and J. Tirole (2011). Identity, morals, and taboos: Beliefs as assets. *The quarterly journal of economics* 126(2), 805–855.
- Benjamini, Y., A. M. Krieger, and D. Yekutieli (2006). Adaptive linear step-up procedures that control the false discovery rate. *Biometrika* 93(3), 491–507.
- Bertrand, M., G. Briscese, M. Grignani, and S. Nassar (2020). How covid-19 is changing americans’ behaviors, expectations, and political views. *The University of Chicago Booth School of Business*. Retrieved September 14, 2022.
- Besley, T. (2006). *Principled Agents? The Political Economy of Good Government*. Oxford University Press.
- Besley, T. and S. Dray (2022). Pandemic responsiveness: Evidence from social distancing and lockdown policy during covid-19. *PloS one* 17(5), e0267611.
- Besley, T. and S. Dray (2024). Trust and state effectiveness: The political economy of compliance. *The Economic Journal*, ueae030.

- Böhm, R., A. Schneider, C. Betsch, and L. Lilleholt (2026). CDC communication undermines trust in vaccines. *Science*.
- Bolsen, T. and J. N. Druckman (2018). Do partisanship and politicization undermine the impact of a scientific consensus message about climate change? *Group Processes & Inter-group Relations* 21(3), 389–402.
- Braithwaite, V. and M. Levi (1998). *Trust and governance*. Russell Sage Foundation.
- Brouwer, N. and J. De Haan (2026). The impact of public knowledge about central bank independence on trust and inflation expectations. *Journal of Economic Behavior & Organization* 241, 107329.
- Budak, C., S. Goel, and J. M. Rao (2016). Fair and balanced? quantifying media bias through crowdsourced content analysis. *Public Opinion Quarterly* 80(S1), 250–271.
- Bursztyn, L., A. Rao, C. Roth, and D. Yanagizawa-Drott (2020). Misinformation during a pandemic. *University of Chicago, Becker Friedman Institute for Economics Working Paper* (2020-44).
- Carril, A. (2017). Randtreat: Stata module to randomly assign treatments uneven treatments and deal with misfits.
- Carson-Chahhoud, K. V., F. Ameer, K. Sayehmiri, K. Hnin, J. E. van Agteren, F. Sayehmiri, M. P. Brinn, A. J. Esterman, A. B. Chang, and B. J. Smith (2017). Mass media interventions for preventing smoking in young people. *Cochrane Database of Systematic Reviews* (6).
- Chetty, R., J. N. Friedman, N. Hendren, M. Stepner, and T. O. I. Team (2020). *The economic impacts of COVID-19: Evidence from a new public database built using private sector data*. Number w27431. National Bureau of Economic Research.
- Clark, C. J., C. Isch, J. A. Everett, and A. Shariff (2023). Even when ideologies align, people distrust politicized institutions.
- Clarke, D., J. P. Romano, and M. Wolf (2020). The romano–wolf multiple-hypothesis correction in stata. *The Stata Journal* 20(4), 812–843.
- Conlon, J. J. (2026). Inattention, information provision, and (failures of) revealed preference. Working paper, Carnegie Mellon University, 21 January 2026.
- Dal Bó, E., P. Dal Bó, and E. Eyster (2018). The demand for bad policy when voters underappreciate equilibrium effects. *The Review of Economic Studies* 85(2), 964–998.
- Darden, M. E. and M. Macis (2024). Trust and health care-seeking behavior. Technical report, National Bureau of Economic Research.
- De Quidt, J., J. Haushofer, and C. Roth (2018). Measuring and bounding experimenter demand. *American Economic Review* 108(11), 3266–3302.

- DellaVigna, S. (2009). Psychology and economics: Evidence from the field. *Journal of Economic literature* 47(2), 315–372.
- Dyer, O. (2020). Covid-19: Black people and other minorities are hardest hit in us.
- Fancourt, D., A. Steptoe, and L. Wright (2020). The cummings effect: politics, trust, and behaviours during the covid-19 pandemic. *The lancet* 396(10249), 464–465.
- Fehr, E., U. Fischbacher, B. v. Rosenbladt, J. Schupp, and G. G. Wagner (2002). A nation-wide laboratory. examining trust and trustworthiness by integrating behavioral experiments into representative surveys. *Journal of Contextual Economics–Schmollers Jahrbuch* (4), 519–542.
- Feldman, L., E. W. Maibach, C. Roser-Renouf, and A. Leiserowitz (2012). Climate on cable: The nature and impact of global warming coverage on fox news, cnn, and msnbc. *The International Journal of Press/Politics* 17(1), 3–31.
- Fuster, A., R. Perez-Truglia, M. Wiederholt, and B. Zafar (2022). Expectations with endogenous information acquisition: An experimental investigation. *Review of Economics and Statistics* 104(5), 1059–1078.
- Gabaix, X. (2014). A sparsity-based model of bounded rationality. *The Quarterly Journal of Economics* 129(4), 1661–1710.
- Gabaix, X. (2019). Behavioral inattention. In *Handbook of behavioral economics: Applications and foundations 1*, Volume 2, pp. 261–343. Elsevier.
- Gaines, B. J., J. H. Kuklinski, P. J. Quirk, B. Peyton, and J. Verkuilen (2007). Same facts, different interpretations: Partisan motivation and opinion on iraq. *The Journal of Politics* 69(4), 957–974.
- Gallup (Last accessed: November 2024). Confidence in u.s. institutions down; average at new low.
- Gaudette, J., S. Hill, T. Kousser, M. Lockhart, M. Romero, and L. Uribe (2026). Can public information campaigns restore trust in American elections across polarized state and party lines? *Election Law Journal: Rules, Politics, and Policy*.
- Gentzkow, M., M. B. Wong, and A. T. Zhang (2025). Ideological bias and trust in information sources. *American Economic Journal: Microeconomics* 17(2), 162–213.
- Grimmelikhuijsen, S. G. and A. J. Meijer (2014). Effects of transparency on the perceived trustworthiness of a government organization: Evidence from an online experiment. *Journal of Public Administration Research and Theory* 24(1), 137–157.
- Groseclose, T. and J. Milyo (2005). A measure of media bias. *The Quarterly Journal of Economics* 120(4), 1191–1237.
- Guiso, L., P. Sapienza, and L. Zingales (2011). Civic capital as the missing link. *Handbook of social economics 1*, 417–480.

- Haaland, I. and C. Roth (2023). Beliefs about racial discrimination and support for pro-black policies. *Review of Economics and Statistics* 105(1), 40–53.
- Haaland, I., C. Roth, and J. Wohlfart (2023). Designing information provision experiments. *Journal of economic literature* 61(1), 3–40.
- Hmielowski, J. D., L. Feldman, T. A. Myers, A. Leiserowitz, and E. Maibach (2014). An attack on science? media use, trust in scientists, and perceptions of global warming. *Public Understanding of Science* 23(7), 866–883.
- Kerwin, J., N. Rostom, and O. Sterck (2024). *Striking the right balance: Why standard balance tests over-reject the null, and how to fix it*. JSTOR.
- Khan, A., S. Nasim, M. Shaukat, and A. Stegmann (2021). Building trust in the state with information: Evidence from urban punjab. *Journal of Public Economics* 202, 104494.
- Kuang, P., M. Weber, and S. Xie (2025, February). Central bank communication with the polarized public. NBER Working Paper 33524, National Bureau of Economic Research.
- León-Ciliotta, G., D. Zejcirovic, and F. Fernandez (2023). Policy-making, trust and the demand for public services: Evidence from a mass sterilization campaign. *American Economic Journal: Economic Policy*.
- Levi, M. (1988). *Of rule and revenue*. Univ of California Press.
- Levi, M. and L. Stoker (2000). Political trust and trustworthiness. *Annual review of political science* 3(1), 475–507.
- Link, S., A. Peichl, O. Pfäuti, C. Roth, and J. Wohlfart (2026). Attention to the macroeconomy. *The Quarterly Journal of Economics*. forthcoming.
- Maćkowiak, B. and M. Wiederholt (2025, April). Rational inattention during an RCT. ECB Working Paper 3054, European Central Bank.
- Martinez-Bravo, M. and C. Sanz (2022). The management of the pandemic and its effects on trust and accountability. *Documentos de Trabajo (CEMFI)* (7), 1.
- Martinez-Bravo, M. and A. Stegmann (2019). In vaccines we trust. *The Effects of the CIA’s Vaccine Ruse on Immunization in Pakistan**(No. wp2017_1713).
- Martinez-Bravo, M. and A. Stegmann (2022). In vaccines we trust? the effects of the cia’s vaccine ruse on immunization in pakistan. *Journal of the European Economic Association* 20(1), 150–186.
- Mayer, R. C., J. H. Davis, and F. D. Schoorman (1995). An integrative model of organizational trust. *Academy of management review* 20(3), 709–734.
- McLaughlin, D. M., J. Mewhirter, and R. Sanders (2021). The belief that politics drive scientific research & its impact on covid-19 risk assessment. *Plos one* 16(4), e0249937.

- Nunnari, S., E. Proto, and A. Rustichini (2024). Cognitive abilities and the demand for bad policy.
- Rafkin, C., A. Shreekumar, and P.-L. Vautrey (2021). When guidance changes: Government stances and public beliefs. *Journal of Public Economics* 196, 104319.
- Ravallion, M., D. Van De Walle, P. Dutta, and R. Murgai (2015). Empowering poor people through public information? lessons from a movie in rural india. *Journal of Public Economics* 132, 13–22.
- Romano, J. P. and M. Wolf (2005). Stepwise multiple testing as formalized data snooping. *Econometrica* 73(4), 1237–1282.
- Sapienza, P., A. Toldra-Simats, and L. Zingales (2013). Understanding trust. *The Economic Journal* 123(573), 1313–1332.
- Sims, C. A. (2003). Implications of rational inattention. *Journal of Monetary Economics* 50(3), 665–690.
- Smith, A. (1759). The theory of moral sentiments.
- Smith, A. (1776). *The wealth of nations [1776]*, Volume 11937. na.
- Sofianos, A. (2022). Self-reported & revealed trust: Experimental evidence. *Journal of Economic Psychology* 88, 102451.
- Stantcheva, S. (2023). How to run surveys: A guide to creating your own identifying variation and revealing the invisible. *Annual Review of Economics* 15(1), 205–234.
- Suhay, E., A. Soni, C. Persico, and D. E. Marcotte (2022). Americans’ trust in government and health behaviors during the covid-19 pandemic. *RSF: The Russell Sage Foundation Journal of the Social Sciences* 8(8), 221–244.
- Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society Series B: Statistical Methodology* 58(1), 267–288.
- Tyler, T. R. (1990). Why people obey the law. *Yale University*.
- Van de Walle, S. and G. Bouckaert (2003). Public service performance and trust in government: The problem of causality. *International journal of public administration* 26(8-9), 891–913.
- Van Scoy, L. J., B. Snyder, E. L. Miller, O. Toyobo, A. Grewel, G. Ha, S. Gillespie, M. Patel, J. Reilly, A. E. Zgierska, et al. (2021). Public anxiety and distrust due to perceived politicization and media sensationalism during early covid-19 media messaging. *Journal of Communication in Healthcare* 14(3), 193–205.
- Wiswall, M. and B. Zafar (2015). Determinants of college major choice: Identification using an information experiment. *The Review of Economic Studies* 82(2), 791–824.

Wright, A. L., K. Sonin, J. Driscoll, and J. Wilson (2020). Poverty and economic dislocation reduce compliance with covid-19 shelter-in-place protocols. *Journal of Economic Behavior & Organization* 180, 544–554.

A Appendix

A1.1 Full Model Derivation and Proofs

This section derives the model in full. We build on the sparsity model of Gabaix (2014, 2019) and depart from it in one place: the citizen holds a belief about how much of an institution’s output is useful to her, and that belief scales the return to attending to anything the institution produces. Everything else — the quadratic loss from inattention, the linear attention cost, the corner solution — is inherited. We state each step as a lemma or proposition and prove it, so that it is clear which conclusions are our own and which are standard.

A1.1.1 Setup

Institutions. Institutions are represented as the canonical benevolent social planner, whose primary goal is to maximize citizens’ well-being by providing evidence-based guidance on behaviors that improve utility and mitigate negative externalities. Building on Braithwaite and Levi (1998) and Tyler (1990), we conceptualize trust as a continuous judgment that influences behavior: higher institutional trust increases individuals’ willingness to comply with recommended behaviors. For institutions to be trustworthy, they must communicate both competence and benevolence (Mayer et al., 1995). We set aside alternative theories framing institutional trust in terms of reciprocity and incentive alignment, focusing instead on how institutions build trust by providing citizens with valuable information that signals competence and integrity.

The decision. A citizen chooses a scalar action $a \in \mathbb{R}$ to maximize a twice continuously differentiable utility $u(a, x)$, where $x = (x_1, \dots, x_n)$ is a vector of attributes of the decision environment and $u_{aa} < 0$. Some attributes are produced by a public institution j : for instance, the air quality index that an EPA tracker reports for the citizen’s zip code. Let $a^*(x)$ denote the fully informed optimum, defined by $u_a(a^*(x), x) = 0$. By the implicit function theorem,

$$a_{x_i} \equiv \frac{\partial a^*}{\partial x_i} = -\frac{u_{ax_i}}{u_{aa}}, \quad (5)$$

which measures how much the optimal action responds to attribute i .

Attention. The citizen does not observe x freely. For each attribute she chooses an attention level $m_i \in [0, 1]$ and perceives

$$x_i^s = m_i x_i + (1 - m_i) x_i^d, \quad (6)$$

a weighted average of the true value and a default x_i^d — the value she would assume if she did not look. With $m_i = 1$ she perceives the truth; with $m_i = 0$ she perceives the default. She then acts on what she perceives, choosing $a^*(x^s)$. Let $\sigma_i^2 \equiv \mathbb{E}[(x_i - x_i^d)^2]$ be the mean squared distance between the truth and the default; when the default is the prior mean, σ_i^2 is simply the variance of x_i . We assume the deviations $x_i - x_i^d$ are mean zero and independent across attributes.

Attention cost. Attention is costly:²⁶

$$\mathcal{C}(m) = \kappa \sum_{i=1}^n c_i m_i, \quad (7)$$

where $\kappa > 0$ is an individual-level cognitive cost and $c_i > 0$ is the search cost of attribute i — how hard it is to find. A direct link lowers c_i ; its absence raises it.

Trust. This is our departure. Not every attribute an institution produces is useful to a given citizen: some of the EPA’s output bears on her decisions and some does not, and she does not know in advance which is which. Let $\omega_i \in \{0, 1\}$ indicate whether attribute i from institution j is useful, in the sense that it enters her decision: the relevant marginal effect is $a_{x_i} \omega_i$, so a useless attribute ($\omega_i = 0$) is one whose value would not change her action even if she knew it. Let $\lambda_j \in [0, 1]$ be the share of j ’s output that is useful to her, and let $\omega_i \mid \lambda_j \sim \text{Bernoulli}(\lambda_j)$ independently across the attributes j produces. She does not know λ_j ; she holds a prior over it with mean

$$\theta_j \equiv \mathbb{E}[\lambda_j] \in [0, 1], \quad (8)$$

and this is what we call trust in j . It is the share of the institution’s work she expects to be useful. The definition is narrower than trust as commonly measured, and deliberately so: it is the component that determines whether she bothers to consult the institution at all, and it is the component our experiments move.

Timing. There are two periods. In period 0 the citizen chooses attention m , given θ_{j0} , then acts. If she attended to an attribute from j she observes whether it was useful, and updates to θ_{j1} . In period 1 she faces the same problem with the updated belief.

A1.1.2 The Loss from Inattention

Lemma 1 (Quadratic loss). *To second order, the expected utility loss from acting on x^s rather than x is*

$$\mathcal{L}(m) = \frac{1}{2} \sum_{i=1}^n (1 - m_i)^2 \Lambda_{ii}, \quad \Lambda_{ii} \equiv \sigma_i^2 a_{x_i}^2 |u_{aa}| = \sigma_i^2 |a_{x_i} u_{aa} a_{x_i}|. \quad (9)$$

Proof. Write $\hat{a} = a^*(x^s)$ for the action taken and $a^* = a^*(x)$ for the fully informed optimum. Expanding $u(\cdot, x)$ around a^* to second order and using $u_a(a^*, x) = 0$,

$$u(a^*, x) - u(\hat{a}, x) \approx -\frac{1}{2} u_{aa} (\hat{a} - a^*)^2 = \frac{1}{2} |u_{aa}| (\hat{a} - a^*)^2.$$

Expanding $a^*(\cdot)$ to first order and using (6),

$$\hat{a} - a^* \approx \sum_{i=1}^n a_{x_i} (x_i^s - x_i) = \sum_{i=1}^n a_{x_i} (1 - m_i) (x_i^d - x_i).$$

²⁶Gabaix (2014) writes the cost as $\kappa \sum_i m_i^\alpha$. We add an attribute-specific search cost c_i and set $\alpha = 1$, the value that makes the solution both sparse and continuous (Tibshirani, 1996).

Squaring, taking expectations, and using that the deviations $x_i^d - x_i$ are mean zero and independent across i so that all cross terms vanish,

$$\mathbb{E}[u(a^*, x) - u(\hat{a}, x)] \approx \frac{1}{2} |u_{aa}| \sum_{i=1}^n a_{x_i}^2 (1 - m_i)^2 \sigma_i^2,$$

which is (9). $\square$

This is the standard result (Gabaix, 2014, Lemma 1). The loss from ignoring attribute i is large when the attribute is variable (σ_i^2), when the action responds to it ($a_{x_i}^2$), and when the action matters ($|u_{aa}|$). Our addition is the next step.

Lemma 2 (Trust scales the loss). *For attributes produced by institution j , the expected loss from inattention, evaluated under the citizen's beliefs about usefulness, is*

$$\mathcal{L}_j(m) = \frac{1}{2} \sum_{i \in j} (1 - m_i)^2 \theta_j \Lambda_{ii}. \quad (10)$$

Proof. For an attribute from j the marginal effect on the action is $a_{x_i} \omega_i$ rather than a_{x_i} . Substituting into the proof of Lemma 1, the i -th term of the loss becomes $\frac{1}{2} |u_{aa}| a_{x_i}^2 \omega_i^2 (1 - m_i)^2 \sigma_i^2$, and the cross terms still vanish because they are multiplied by $\mathbb{E}[(x_i^d - x_i)(x_k^d - x_k)] = 0$. Since $\omega_i \in \{0, 1\}$, $\omega_i^2 = \omega_i$. The citizen chooses m before observing ω_i and ω_i is independent of $x_i - x_i^d$, so taking expectations over her beliefs,

$$\mathbb{E}[\omega_i] = \mathbb{E}[\mathbb{E}[\omega_i | \lambda_j]] = \mathbb{E}[\lambda_j] = \theta_j$$

by the law of iterated expectations and (8). Hence the i -th term is $\frac{1}{2} (1 - m_i)^2 \theta_j \Lambda_{ii}$. $\square$

The intuition is that an attribute the citizen expects to be useless is costless to ignore. Trust therefore enters as a multiplier on the return to attention for everything j produces, not as a shift in any single attribute's relevance. Setting $\theta_j = 1$ recovers Gabaix (2014) exactly.

A1.1.3 Optimal Attention

The citizen chooses $m \in [0, 1]^n$ to minimize expected loss plus attention cost:

$$\min_{m \in [0, 1]^n} \frac{1}{2} \sum_{i=1}^n (1 - m_i)^2 \theta_j \Lambda_{ii} + \kappa \sum_{i=1}^n c_i m_i. \quad (11)$$

Proposition 1 (Optimal attention). *The solution to (11) is, for each attribute i ,*

$$m_i^* = A\left(\frac{\theta_j \Lambda_{ii}}{\kappa c_i}\right), \quad A(z) \equiv \max\{0, 1 - 1/z\}, \quad (12)$$

with the convention $A(0) = 0$.

Proof. The objective is additively separable across i , so each m_i can be chosen independently. Fix i and write $f(m_i) = \frac{1}{2}(1 - m_i)^2\theta_j\Lambda_{ii} + \kappa c_i m_i$. If $\theta_j\Lambda_{ii} = 0$ then f is linear with slope $\kappa c_i > 0$ and the minimizer is $m_i = 0$, consistent with the convention $A(0) = 0$. Otherwise $f''(m_i) = \theta_j\Lambda_{ii} > 0$, so f is strictly convex on $[0, 1]$ and has a unique minimizer. Its derivative is

$$f'(m_i) = -(1 - m_i)\theta_j\Lambda_{ii} + \kappa c_i.$$

The unconstrained stationary point is $\tilde{m}_i = 1 - \kappa c_i/(\theta_j\Lambda_{ii})$, which is strictly below 1 since $\kappa c_i > 0$. Two cases. If $\tilde{m}_i > 0$ it lies in $(0, 1)$ and is the constrained minimizer. If $\tilde{m}_i \leq 0$, then $f'(0) = -\theta_j\Lambda_{ii} + \kappa c_i \geq 0$, and by convexity $f' \geq 0$ on all of $[0, 1]$, so the minimizer is the corner $m_i = 0$. Together, $m_i^* = \max\{0, \tilde{m}_i\}$, which is (12) with $z = \theta_j\Lambda_{ii}/(\kappa c_i)$. $\square$

Corollary 1 (Threshold). *Attribute i receives positive attention if and only if*

$$\theta_j > \bar{\theta}_i \equiv \frac{\kappa c_i}{\Lambda_{ii}} = \frac{\kappa c_i}{\sigma_i^2 |a_{x_i} u_{aa} a_{x_i}|}. \quad (13)$$

Proof. From the proof of Proposition 1, $m_i^* > 0$ if and only if $\tilde{m}_i > 0$, i.e. $\kappa c_i < \theta_j\Lambda_{ii}$. $\square$

The corner is the feature our design exploits. Three distinct conditions each produce *zero* attention rather than merely less of it: low trust (θ_j small), a high search cost (c_i large), and information of little decision relevance (small a_{x_i} , hence small Λ_{ii} and a high threshold). Our first experiment and the T1–T2 contrast of the second lower c_i ; the T1–T3 contrast varies a_{x_i} ; and the mechanism tables in Section 4.3.1 concern θ_j .

A1.1.4 Interpretation of Terms

Four terms govern optimal attention.

θ_j is trust: the share of institution j 's work the citizen expects to be useful. By Lemma 2 it scales the entire benefit of attending, so it acts as a gatekeeper on the whole of j 's output rather than on any single attribute.

σ_i^2 is the variability of attribute i around the default — how much there is to learn by looking. A parent in an area with consistently clean air has little reason to check the AQI regularly, whereas one in an area with volatile pollution levels faces genuine health risk from ignoring it.

a_{x_i} is how much a change in x_i shifts the optimal action. Information that is consequential for decisions receives more attention.

u_{aa} is how much the individual cares about the quality of the action — the stakes.

Together, an individual pays more attention to attributes that are more variable, more consequential, less costly to acquire, and produced by an institution she expects to be useful.

A1.1.5 Updating and the Vicious Cycle

Having chosen m^* , the citizen acts, $a^* = \arg \max_a u(a, x^s)$, and then learns from what she consulted.

Assumption 1 (Beta prior). *The citizen's prior over λ_j is Beta(α_j, β_j) with $\alpha_j, \beta_j > 0$. Its mean is $\theta_{j0} = \alpha_j / (\alpha_j + \beta_j)$ and we write $n_j \equiv \alpha_j + \beta_j$ for its precision, which can be read as the amount of prior experience she has with j .*

Proposition 2 (Updating). *If $m_i^* > 0$ for some attribute i from j , the citizen observes $\omega_i \sim \text{Bernoulli}(\lambda_j)$ and her posterior mean is*

$$\theta_{j1} = \theta_{j0} + \frac{1}{n_j + 1} (\omega_i - \theta_{j0}). \quad (14)$$

If $m_i^ = 0$ for every attribute from j , she observes nothing and $\theta_{j1} = \theta_{j0}$.*

Proof. The Beta distribution is conjugate to the Bernoulli: a Beta(α, β) prior combined with one Bernoulli observation ω gives a Beta($\alpha + \omega, \beta + 1 - \omega$) posterior. Its mean is

$$\theta_{j1} = \frac{\alpha_j + \omega_i}{\alpha_j + \beta_j + 1} = \frac{n_j \theta_{j0} + \omega_i}{n_j + 1} = \theta_{j0} + \frac{\omega_i - \theta_{j0}}{n_j + 1},$$

using $\alpha_j = n_j \theta_{j0}$. If nothing is observed, the posterior equals the prior. $\square$

Trust moves toward the usefulness the citizen experienced, by an amount that shrinks with her prior experience: a first encounter moves beliefs a lot, the hundredth hardly at all. Two remarks on what does *not* move trust.

Remark 1 (No consultation, no update). The update in (14) is conditional on having attended. A citizen with $m_i^* = 0$ observes no ω_i and her belief is unchanged, whatever the true λ_j .

Remark 2 (Low-value information: less attention, and a weaker signal). An attribute of low decision relevance — small but positive a_{x_i} , such as the city in which a regional office is located — affects both stages. In the attention stage, small a_{x_i} means small Λ_{ii} and hence a high threshold $\bar{\theta}_i$, so fewer citizens consult it than would consult a high-value attribute at the same search cost; with the search cost near zero, as when a link is supplied, some still do. In the updating stage, what a citizen learns from consulting it is weaker. The usefulness she can experience is bounded by how much the attribute could have changed her action, so the probability that it is found useful is $\lambda_j \rho_i$ rather than λ_j , with $\rho_i \in (0, 1]$ increasing in Λ_{ii} and equal to one for the most relevant attribute. The observation is then less informative about λ_j , and the expected revision of θ_j from consulting attribute i is smaller in magnitude, approximately ρ_i times the revision from a fully relevant attribute. The model therefore predicts that T3 in Study II should show positive but lower engagement than T1 — 54 against 66 percent opened the link — and a smaller effect on trust. We treat ρ_i as a reduced-form description of this attenuation rather than deriving it.

Remark 3 (Relevance and distortion). Write usefulness as $\omega_i = r_i (1 - d_i)$, where $r_i \in \{0, 1\}$ indicates that attribute i bears on the citizen's decision and $d_i \in \{0, 1\}$ that it is distorted, so

that even a relevant attribute is useless if it is reported in a way that does not track the truth. Let $\lambda_j = \lambda_j^r (1 - \lambda_j^d)$, with λ_j^r the share of j 's output that is relevant and λ_j^d the share that is distorted, and let the citizen hold independent priors over the two. Perceived politicization is her belief about λ_j^d . The prompt in T1 and T2 reveals $r_i = 1$ before any consultation, so consulting the page is uninformative about λ_j^r and informative only about λ_j^d : a number that matches what the citizen can verify is a draw of $d_i = 0$. The T3 prompt reveals a low r_i , and by Remark 2 the page then carries little information about either share. Under this decomposition the T1–T3 contrast moves beliefs about distortion, which is what the independence question elicits, more than beliefs about relevance, which is closer to what the benevolence question elicits, and the update is larger for citizens who can check the number. This is consistent with the pattern in Tables 2 and 3, but we offer it as an interpretation rather than as a prediction the design was built to test.

Proposition 3 (The vicious cycle). *Suppose $\theta_{j0} \leq \bar{\theta}_i$ for every attribute i that j produces. Then $\theta_{jt} = \theta_{j0}$ for all $t \geq 0$: trust in j never changes.*

Proof. By Corollary 1, $m_i^* = 0$ for every attribute from j in period 0. By Proposition 2 the citizen observes nothing and $\theta_{j1} = \theta_{j0}$. She therefore faces an identical problem in period 1, and the argument repeats. $\square$

The cycle is self-sealing rather than self-confirming. A citizen below the threshold never acquires the evidence that might revise her belief, so mistrust persists regardless of whether it was warranted; the same mechanism sustains high trust above the threshold. An intervention that lowers c_i lowers $\bar{\theta}_i$ and can bring a citizen over it, after which (14) operates and trust can move. The size of that movement is what the next result decomposes.

A1.1.6 The Intention-to-Treat Effect on Trust

Let $D \in \{0, 1\}$ indicate that the citizen consulted the institution ($m_i^* > 0$ for some i from j), and let $\Delta\theta_j \equiv \theta_{j1} - \theta_{j0}$. Compare a treated arm T and a comparison arm C , and write $\pi_T = \Pr(D = 1 \mid T)$ and $\pi_C = \Pr(D = 1 \mid C)$ for the shares consulting.

Proposition 4 (ITT decomposition). *The intention-to-treat effect of the treatment on trust is*

$$ITT_{trust} = \pi_T \mathbb{E}[\Delta\theta_j \mid D = 1, T] - \pi_C \mathbb{E}[\Delta\theta_j \mid D = 1, C]. \quad (15)$$

If, in addition, consulting yields the same expected update in both arms, $\mathbb{E}[\Delta\theta_j \mid D = 1, T] = \mathbb{E}[\Delta\theta_j \mid D = 1, C] \equiv \delta$, then

$$ITT_{trust} = \underbrace{(\pi_T - \pi_C)}_{\text{change in share consulting}} \times \underbrace{\delta}_{\text{update among consumers}}. \quad (16)$$

Proof. By the law of iterated expectations, for either arm $g \in \{T, C\}$,

$$\mathbb{E}[\Delta\theta_j \mid g] = \pi_g \mathbb{E}[\Delta\theta_j \mid D = 1, g] + (1 - \pi_g) \mathbb{E}[\Delta\theta_j \mid D = 0, g] = \pi_g \mathbb{E}[\Delta\theta_j \mid D = 1, g],$$

because by Remark 1 $\Delta\theta_j = 0$ whenever $D = 0$. Subtracting the C expression from the T expression gives (15). Under the equal-update condition both conditional expectations equal δ , and (15) factors into (16). $\square$

The equal-update condition is what distinguishes the two contrasts in Study II. T1 and T2 give access to the *same* EPA page; a respondent who reaches it experiences the same ω_i whichever arm she is in, so δ is common and (16) applies. Their trust effects then differ only by the difference in the share who reach the page, which can be small even when δ is large — the point made in Section 4.3. T1 and T3 give access to *different* pages, and by Remark 2 the update from the T3 page is smaller, $\delta_{T3} = \rho_{T3} \delta_{T1}$ with $\rho_{T3} < 1$, so the equal-update condition fails and (15) applies: the T1–T3 trust effect is $\pi_{T1} \delta_{T1} - \pi_{T3} \delta_{T3}$, positive both because more T1 respondents consult and because each learns more. The comparison among those who clicked in Table 3 estimates the second component, $\delta_{T1} - \delta_{T3}$, under the further assumption that respondents who click in T1 and those who click in T3 start from the same θ_{j0} ; since clicking is a choice, that assumption is not guaranteed, which is why we present those estimates as descriptive.

A1.1.7 Example: Quadratic Utility

The general expressions simplify under quadratic utility, and this is the case used in the numerical example that follows.

Proposition 5 (Quadratic utility). *Let $u(a, x) = -\frac{1}{2}(a - \sum_{i=1}^n b_i x_i)^2$. Then $a_{x_i} = b_i$, $u_{aa} = -1$, $\Lambda_{ii} = \sigma_i^2 b_i^2$, and*

$$m_i^* = \max \left\{ 0, 1 - \frac{\kappa c_i}{\theta_j \sigma_i^2 b_i^2} \right\}, \quad (17)$$

with positive attention if and only if $\kappa c_i < \theta_j \sigma_i^2 b_i^2$.

Proof. Differentiating, $u_a = -(a - \sum_i b_i x_i)$, so $u_{aa} = -1$ and $u_{ax_i} = b_i$. By (5), $a_{x_i} = -u_{ax_i}/u_{aa} = b_i$, and the fully informed optimum is $a^*(x) = \sum_i b_i x_i$. Substituting into (9), $\Lambda_{ii} = \sigma_i^2 b_i^2 \cdot 1$. Proposition 1 then gives (17), and Corollary 1 gives the condition. $\square$

The number of attributes that can receive positive attention is therefore bounded: at most $\sum_i \mathbf{1}[\kappa c_i < \theta_j \sigma_i^2 b_i^2]$ of them. Higher cognitive costs κ , or lower trust θ_j , shrink that set.

A1.1.8 A Numerical Example

The following example makes the mechanism concrete and shows how the two experiments map onto the parameters of equation (1). It is illustrative rather than calibrated: the parameters are chosen to be transparent, not estimated.

Under quadratic utility, equation (12) specialises to $m_i^* = \max\{0, 1 - \kappa c_i / (\theta_j \sigma_i^2 b_i^2)\}$, so attention is strictly positive if and only if trust clears the threshold in equation (13), $\theta_j > \bar{\theta}_i = \kappa c_i / (\sigma_i^2 b_i^2)$.

Set the cognitive cost to $\kappa = 1$ and consider two pieces of information. *High-value* information — the local air quality index in Study II, the current COVID-19 death count in Study I — is variable and consequential for the decision at hand, $\sigma^2 b^2 = 1.00$. *Low-value* information — the city hosting the EPA regional office — is accurate but of little relevance to any decision, $\sigma^2 b^2 = 0.55$. Consider also two search costs: a direct link to the institution's

page, $c = 0.30$, and no link, $c = 0.85$. Finally, consider a trusting citizen, $\theta_j = 1$, and a mistrusting one, $\theta_j = 0.5$. Table A1 reports the resulting optimal attention.

Table A1: Optimal attention m^* in the numerical example.

		High-value info. $\sigma^2 b^2 = 1.00$	Low-value info. $\sigma^2 b^2 = 0.55$
Trusting citizen ($\theta_j = 1$)	Link ($c = 0.30$)	0.70	0.45
	No link ($c = 0.85$)	0.15	0
Mistrusting citizen ($\theta_j = 0.5$)	Link ($c = 0.30$)	0.40	0
	No link ($c = 0.85$)	0	0

Notes: Optimal attention $m^* = \max\{0, 1 - \kappa c / (\theta_j \sigma^2 b^2)\}$ with $\kappa = 1$. The threshold level of trust above which attention is positive, $\bar{\theta} = \kappa c / (\sigma^2 b^2)$, is 0.30 for high-value information with a link, 0.85 without a link, 0.55 for low-value information with a link, and 1.55 for low-value information without a link — above one, so no citizen attends.

Three features of the table map directly onto the experimental arms. First, reading down the trusting citizen’s rows in the first column reproduces the search-cost comparison, T1 versus T2 in Study II and treatment versus control in Study I: removing the link cuts attention from 0.70 to 0.15, because the threshold rises from $\bar{\theta} = 0.30$ to $\bar{\theta} = 0.85$. Second, reading across that citizen’s first row reproduces the information-value comparison, T1 versus T3: holding the search cost fixed, low-value information yields 0.45 rather than 0.70, and the ordering $m^*(T1) > m^*(T3) > m^*(T2)$ matches the ordering of observed engagement across the three arms of Study II: 66 percent of T1 respondents opened the link against 54 percent in T3, while T2, with no link to open, answered correctly least often of all (17 percent at the ± 5 window).

Third, and central to the interpretation of our heterogeneity results, the mistrusting citizen attends only in the single cell where the information is both valuable and cheap to obtain. When the link is withdrawn, her attention is zero — not merely lower — because $\bar{\theta} = 0.85$ exceeds her $\theta_j = 0.5$. The model therefore predicts that in the high-search-cost condition the only respondents who acquire the information are those with a low enough perceived cost: the trusting, and those who already follow the news. Section A2 shows that this is what the data look like in both studies. It also predicts that lowering the search cost is what brings mistrusting citizens over the threshold (0 to 0.40), but only when the information is valuable: with low-value information the mistrusting citizen stays at zero even when the link is provided, and those who do cross the threshold learn less from what they find, which is why T3 moves trust less than T1 does.

A1.1.9 Model Limitations

Our model focuses on individual-level attention costs and information value as determinants of trust. It abstracts from the macro-level forces behind the sharp aggregate decline in institutional trust — including elite political attacks, partisan media narratives, and the psychological shock of an unprecedented public health crisis. These forces likely operate

through channels our model does not formalize: they may raise the perceived cost of engaging with institutional information (higher c_i for opposing partisans), reduce the credibility weight placed on such information, or shift priors about institutional competence in ways unrelated to individual attention allocation. We view these macro-level forces and our demand-side attention frictions as complementary rather than competing mechanisms: the former helps explain why trust declined so sharply, while our framework explains why it may be recoverable through targeted informational interventions.

A2 Prior Beliefs, Trust, and Who Pays the Search Cost

The numerical example in Appendix A1.1.8 predicts that when the search cost is high, attention is positive only for citizens whose perceived cost of engaging with the institution is low enough to clear the threshold in equation (13) — the trusting, and those who already follow the news. This subsection reports the evidence for that prediction in both studies. Two caveats apply throughout. These are correlations between pre-treatment characteristics and outcomes, not causal estimates: news consumption and trust in science are chosen, not assigned. And the two studies measure the prior differently. Study I has a measure of trust in the scientific community collected in wave 1 (April 2020), two months before the experiment ran in wave 6, and so genuinely pre-treatment; Study II has no pre-treatment measure of trust in science, so we use the breadth and slant of news consumption alone.

A2.0.1 Study I

Table A2 reports the results. Respondents who consume more news are substantially more likely to answer the death-toll questions correctly. Following U.S. news for at least thirty minutes a day is associated with a 12.4 percentage point higher probability of a correct answer without controls and 9.7 points with the full control set (cols. 1–2), and following COVID-19 news more closely raises it by 5.9 to 8.1 points per point on the five-point scale (cols. 3–4). Both are large relative to the 39 percent unconditional accuracy rate, and both survive the full set of demographic, local-pandemic and news-slant controls. The slant of news consumption, by contrast, predicts nothing once weights are applied: neither the Republican-leaning nor the Democratic-leaning indicator is significant in any specification. What matters is how much news a respondent consumes, not which side it comes from.

Pre-treatment trust in the scientific community predicts accuracy in the same direction. Respondents who reported no more than “some confidence” in the scientific community in April — 36 percent of the experimental sample — answered correctly 29.6 percent of the time, against 44.4 percent among the more trusting, a raw gap of 14.8 points. The gap is 12.2 points conditional on treatment status alone and 7.2 points with the full control set (cols. 5–6). When news consumption and trust are entered together (col. 7), the news measures retain their size and significance while the trust coefficient falls to -4.5 points and is no longer significant, which is what one would expect if trust in science and habitual news consumption are two measurements of the same underlying propensity to engage with institutional information.

The most direct test concerns the control group, where no link was offered and the

search cost was therefore high. Table A3 restricts attention to those respondents. Trust in science predicts how long they spent on the question: low-trust respondents spent 18.8 fewer seconds without controls and 13.9 fewer with them, against a survey-weighted control-group mean of 52.1 seconds, and were 9.0 points less likely to answer correctly (7.0 with controls). The raw, unweighted contrast is stark: 43.6 seconds and 23.3 percent accuracy among the low-trust, against 58.1 seconds and 38.1 percent among the rest. Heavy news consumers move in the opposite direction, spending 14.7 more seconds and answering correctly 14.3 points more often. Consumers of Republican-leaning news spent 25.1 fewer seconds with controls and were 12.8 points less likely to fall in the top quartile of time spent; consumers of Democratic-leaning news, entered in the same regression, show no such deficit (-2.9 points, not significant), so the slant effect on searching is specific to the right.

The time result is specific to the arm in which searching was costly, and this is the one place where we can show that rather than assert it. Interacting treatment status with pre-treatment trust, the low-trust deficit in seconds spent is -19.3 (s.e. 7.9) in the control group and $+12.5$ (s.e. 11.0, not significant) among the treated, an interaction of 31.8 with a standard error of 13.7 ($p = 0.022$). When the institution’s information is a click away, the gap in time spent between trusting and mistrusting respondents closes entirely. This is the pattern in equation (13): when the cost of searching is high it is paid by those for whom it is perceived as low, and lowering that cost is what brings the others in. The accuracy gap behaves differently: low trust predicts a 9.1 point deficit in the control group and a 15.4 point deficit among the treated, and the two cannot be distinguished ($p = 0.32$), so we do not claim that the knowledge gap itself is concentrated where search costs are high.

Time spent is in turn strongly associated with accuracy. Respondents in the top quartile of time on the death questions are 25.1 to 27.3 percentage points more likely to answer correctly (Table A5, cols. 1–2). The treatment itself raises time spent by roughly 14.5 seconds, consistent with the interpretation of the link as a reduction in c_i that raises m_i^* .

Finally, the treatment effect itself is concentrated among the more trusting. Splitting on pre-treatment trust, the effect of the link on accuracy is 11.9 points and significant at the 1 percent level among respondents with higher trust in science ($n = 725$), and 6.2 points and not significant among those with lower trust ($n = 415$). We report this as suggestive only: the interaction term is -6.3 points with a standard error of 6.3 and cannot be distinguished from zero, so the study is not powered to establish that the difference is real.

A2.0.2 Study II

Study II reproduces the news-consumption result with an independent sample, a different institution and a different factual question. Table A4 reports it. Each additional news source a respondent reports is associated with a 1.6 percentage point higher probability of answering correctly without controls and 1.2 points with them (cols. 1–2); consuming three or more sources — 48 percent of the sample — raises it by 6.9 points (4.7 with controls, cols. 3–4). Unlike in Study I, slant does matter here: consumers of Republican-leaning news are 6.7 points less likely to answer correctly (cols. 5–6), which is unsurprising for a question about the EPA, an agency on which partisan priors are sharply divided in a way they were not for the COVID-19 death count in June 2020.

Restricting to T2, the arm that received the air quality question with no link, each

additional news source is associated with a 2.1 point higher probability of a correct answer without controls and 2.3 points with them (cols. 7–8), significant at the 5 percent level in both specifications despite the smaller sample. We caution against reading this as evidence that the gradient is steeper where the search cost is higher. The within-arm gradient is essentially identical in T1, which asked the same air quality question with a link (2.06 points per source, $p = 0.03$), and flat in T3, which asked about the location of the regional office (0.58 points, $p = 0.58$); a joint test of the two interaction terms cannot reject equality across arms ($p = 0.51$ without controls, $p = 0.25$ with them). What separates the arms is the content of the question rather than the cost of answering it: following the news helps with the current air quality index and is of no use for the city hosting an administrative office. This is consistent with the $\sigma_i^2 b_i^2$ term of the model rather than with its c_i term, and it is why the pooled coefficient in columns (1)–(2) is smaller than the coefficient in either air quality arm.

Study II also observes the two intermediate margins directly. Clicking the link is associated with a 39 to 40 percentage point higher probability of a correct answer, and being in the top quartile of time spent on the question with 18.5 points; entered jointly, both remain significant, at 37.0 and 9.6 points (Table A5, cols. 3–6). Time is measured the same way in both studies, as the top quartile of seconds spent on the treatment question. These are the attention margins m_i^* of the model, and they are the strongest single predictors of accuracy anywhere in either study. Among respondents offered the air quality link, those who consume Democratic-leaning news are 10.4 to 12.3 points more likely to click, while news breadth predicts neither clicking nor time spent — consistent with the reading that breadth captures a stock of prior knowledge that pays off mainly where no link is available, rather than a greater propensity to search when one is.

Table A2: News consumption, pre-treatment trust in science, and knowledge accuracy (Study I).

VARIABLES	(1)	(2)	(3)	(4)	(5)	(6)	(7)
				Correct answer			
30+ min/day of U.S. news	0.124*** (0.0272)	0.0969*** (0.0245)					0.0663** (0.0276)
Follows COVID-19 news (1–5)			0.0807*** (0.0182)	0.0585*** (0.0174)			0.0464** (0.0186)
Low trust in science (w1)					-0.122*** (0.0349)	-0.0718** (0.0303)	-0.0450 (0.0302)
Treated	0.0999*** (0.0332)	0.125*** (0.0286)	0.101*** (0.0331)	0.124*** (0.0285)	0.0993*** (0.0315)	0.124*** (0.0282)	0.127*** (0.0269)
Observations	1,147	1,147	1,147	1,147	1,147	1,147	1,147
Demographics	No	Yes	No	Yes	No	Yes	Yes
Covid shock	No	Yes	No	Yes	No	Yes	Yes
News slant	No	Yes	No	Yes	No	Yes	Yes
Mean dep. var.	0.391	0.391	0.391	0.391	0.391	0.391	0.391

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

Survey-weighted OLS regressions on the wave 6 experimental sample. The outcome is whether the respondent correctly identified both the state and the national COVID-19 death toll. *Low trust in science* equals one for respondents reporting at most “some confidence” in the scientific community in wave 1 (April 2020), two months before the experiment, and is therefore measured pre-treatment. *News slant* denotes the Republican- and Democratic-leaning news indicators used throughout the paper; neither is significant in any column. The control variables are those listed in Appendix A4.0.2. Produced by CDC/Codes/5.news_trust_correct.do.

Table A3: Who pays the search cost when it is high? Control group only (Study I).

VARIABLES	(1)	(2)	(3)	(4)	(5)	(6)
	Seconds on question		Long time on question		Correct answer	
Low trust in science (w1)	-18.82** (7.920)	-13.87** (5.962)			-0.0903** (0.0440)	-0.0696** (0.0352)
30+ min/day of U.S. news	14.66* (8.107)	1.662 (7.488)			0.143*** (0.0374)	0.0897** (0.0349)
Rep-leaning news			-0.101* (0.0575)	-0.128*** (0.0492)		
Dem-leaning news			-0.0173 (0.0570)	-0.0294 (0.0493)		
Observations	1,175	1,175	1,175	1,175	1,168	1,168
Demographics	No	Yes	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes	No	Yes
Mean dep. var.	52.13	52.13	0.250	0.250	0.280	0.280

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

Survey-weighted OLS regressions restricted to the control group, which received the death-toll questions without a link and therefore faced the higher search cost. *Long time on question* equals one for the top quartile of seconds spent ($\geq$ 75th percentile). Low trust and news volume each come from a separate regression of the outcome on that variable; the Republican- and Democratic-leaning news indicators are entered jointly, with no slanted news as the omitted category. The mean of the dependent variable is the survey-weighted control-group mean. Raw, unweighted means in the control group: respondents with low trust in science spent 43.6 seconds and answered correctly 23.3 percent of the time, against 58.1 seconds and 38.1 percent for the rest. Produced by CDC/Codes/5.news_trust_correct.do.

Table A4: News consumption and knowledge accuracy (Study II).

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
	All arms				T2 only (no link)			
VARIABLES	Correct response							
N. of news sources	0.0160*** (0.00560)	0.0122** (0.00590)					0.0208** (0.00870)	0.0225** (0.00930)
Three or more news sources			0.0686*** (0.0255)	0.0474* (0.0271)				
Rep-leaning news					-0.0667** (0.0338)	-0.0681* (0.0351)		
Observations	1,400	1,400	1,400	1,400	1,400	1,400	535	535
Treatment arms	Yes	Yes	Yes	Yes	Yes	Yes	—	—
Demographics	No	Yes	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.441	0.441	0.441	0.441	0.441	0.441	0.271	0.271

Notes: Standard errors in parentheses, robust to heteroskedasticity. *** p<0.01, ** p<0.05, * p<0.1.

OLS regressions. The outcome is whether the respondent answered the factual question correctly, scored against the corrected AQI benchmark described in Section 4.2. *N. of news sources* counts how many of the twelve listed outlets the respondent reports using. Columns (1)–(6) include the treatment arm indicators; columns (7)–(8) are restricted to T2, which received the air quality question without a link and therefore faced the higher search cost. Produced by TESS/Codes/EPA/9.news_trust_correct.do.

Table A5: Attention and knowledge accuracy in both studies.

	(1)	(2)	(3)	(4)	(5)	(6)
	Study I		Study II			
VARIABLES	Correct answer					
Long time on question	0.273*** (0.0385)	0.251*** (0.0355)				
Clicked link			0.404*** (0.0317)	0.389*** (0.0332)		0.370*** (0.0343)
Long time on question					0.185*** (0.0370)	0.0960*** (0.0368)
Observations	1,132	1,132	865	865	865	865
Demographics	No	Yes	No	Yes	Yes	Yes
News	No	Yes	No	Yes	Yes	Yes

Notes: Standard errors in parentheses; survey-weighted in cols. (1)–(2), robust in cols. (3)–(6). *** p<0.01, ** p<0.05, * p<0.1.

All columns control for treatment status. Columns (1)–(2) use the wave 6 experimental sample of Study I; *Long time on question* equals one for the top quartile of seconds spent on the death-toll questions. Columns (3)–(6) use the two arms of Study II that were offered a link (T1 and T3), the only arms in which clicking and time on the question are recorded; there *Long time on question* is built the same way, as the top quartile of seconds spent on the treatment question (≥ 60 seconds). These are associations between two endogenous choices and the outcome, not causal estimates. Produced by CDC/Codes/5.news_trust_correct.do and TESS/Codes/EPA/9.news_trust_correct.do.

A3 Multiple-Hypothesis Adjustments

The main tables report unadjusted p-values for several outcomes at a time. This section reports two adjustments for each family of outcomes: the sharpened false-discovery-rate q-values by Anderson (2008), and the step-down family-wise adjusted p-values of Romano and Wolf (2005), computed by bootstrap with 2,000 replications (Clarke et al., 2020). A q-value

of 0.05 means that, among outcomes declared significant at that threshold, an expected 5 percent are false discoveries; a Romano-Wolf p-value controls the probability of any false rejection within the family. In Study I the families are the four knowledge and trust outcomes of Table 1 and the four behavioural outcomes of Table A16; the bootstrap resamples primary sampling units and applies the survey weights. As a stricter alternative we also pool all eight Study I outcomes into a single family. In Study II the family is the three pre-registered binary trust outcomes of Table 2, separately for the T1–T3 and T2–T3 contrasts; a pooled family of six is also reported. Knowledge accuracy in Study II is significant at any conventional level under any adjustment ($p < 0.001$) and is omitted.

Table A6: Multiple-hypothesis adjustments, Study I (Tables 1 and A16).

Controls	Outcome	Coef. (s.e.)	Unadj. p	Sharpened q (4)	Sharpened q (8)	Romano-Wolf p (4)
<i>Family A: knowledge and trust</i>						
No	Correct death est. (June)	0.101 (0.032)	0.002	0.009	0.018	0.006
No	Correct US death rank (Oct)	0.041 (0.034)	0.229	0.107	0.131	0.230
No	Incr. trust in CDC	0.064 (0.033)	0.052	0.078	0.091	0.032
No	CDC not infl. by politics	0.055 (0.030)	0.072	0.078	0.091	0.050
Yes	Correct death est. (June)	0.125 (0.029)	0.000	0.001	0.001	0.008
Yes	Correct US death rank (Oct)	0.054 (0.031)	0.080	0.051	0.061	0.158
Yes	Incr. trust in CDC	0.062 (0.029)	0.036	0.038	0.053	0.060
Yes	CDC not infl. by politics	0.060 (0.027)	0.028	0.038	0.053	0.060
<i>Family B: behaviours</i>						
No	Tested for Covid-19	0.080 (0.037)	0.031	0.067	0.079	0.294
No	Adopts no prev. behav.	−0.024 (0.011)	0.027	0.067	0.079	0.294
No	% prev. behav. adopted	0.002 (0.022)	0.916	0.845	0.440	0.745
No	Covers mouth in public	0.022 (0.029)	0.457	0.439	0.244	0.294
Yes	Tested for Covid-19	0.066 (0.031)	0.035	0.057	0.053	0.197
Yes	Adopts no prev. behav.	−0.023 (0.009)	0.013	0.057	0.049	0.197
Yes	% prev. behav. adopted	−0.001 (0.020)	0.961	0.925	0.317	0.953
Yes	Covers mouth in public	0.019 (0.023)	0.402	0.366	0.130	0.313

Notes: Coefficients, standard errors and unadjusted p-values are the survey-design estimates of Tables 1 and A16 (survey weights, primary sampling units, strata). *Sharpened q (4)*: Benjamini et al. (2006) q-values within the family of four outcomes for the given control set. *Sharpened q (8)*: the same, pooling both families. *Romano-Wolf p (4)*: step-down adjusted p-values within the family of four, from 2,000 survey-weighted bootstrap replications that resample primary sampling units (seed 20260915). In family B the bootstrap distribution of the t-statistic for *Adopts no preventive behaviour*, whose control mean is 2 percent, is wide, and it is this outcome that sets the family-wise adjustment for the whole family. Produced by CDC/Codes/8.mht.cdc.do.

Table A7: Multiple-hypothesis adjustments, Study II (Table 2, cols 3–8).


Controls	Outcome	Coef. (s.e.)	T1 (low search costs, high-value information) vs T3			
			Unadj. p	Sharpened q (3)	Sharpened q (6)	Romano-Wolf p (3)
No	EPA improves lives	0.050 (0.030)	0.093	0.055	0.142	0.095
No	EPA independent	0.062 (0.025)	0.013	0.042	0.087	0.045
No	Support for EPA	0.063 (0.030)	0.035	0.042	0.095	0.066
Yes	EPA improves lives	0.023 (0.029)	0.415	0.179	0.497	0.434
Yes	EPA independent	0.046 (0.025)	0.063	0.179	0.435	0.170
Yes	Support for EPA	0.048 (0.029)	0.101	0.179	0.435	0.178
T2 (high search costs, high-value information) vs T3						
No	EPA improves lives	0.018 (0.029)	0.543	0.382	0.284	0.552
No	EPA independent	0.034 (0.024)	0.171	0.382	0.173	0.399
No	Support for EPA	0.039 (0.030)	0.184	0.382	0.173	0.399
Yes	EPA improves lives	0.012 (0.028)	0.677	0.795	0.795	0.679
Yes	EPA independent	0.025 (0.024)	0.295	0.795	0.435	0.564
Yes	Support for EPA	0.033 (0.029)	0.249	0.795	0.435	0.564

Notes: Coefficients, classical standard errors and unadjusted p-values reproduce Table 2, cols (3)–(8) ($N = 1,626; 1,613; 1,580$). *Sharpened q (3)*: Benjamini et al. (2006) q-values within the family of three trust outcomes for the given contrast and control set. *Sharpened q (6)*: the same, pooling the T1 and T2 contrasts into one family of six. *Romano-Wolf p (3)*: step-down adjusted p-values within the family of three, 2,000 bootstrap replications, seed 20260915. Produced by TESS/Codes/EPA/15.mht.socialmedia.do.

A4 Study I

A4.0.1 Design

Figure A1: Randomized treatment



Notes: The figure shows the webpage that treated respondents would land on if they clicked on the link.

A4.0.2 Controls

Economic shocks. To measure direct economic shocks, we asked all respondents in the last wave of the survey to report their (and their spouse, if present) monthly gross income between February and October²⁷. Furthermore, we also asked respondents (and their spouses, if present) about their monthly additional sources of income, the number of hours they worked each month, and whether they had received any financial support from the government or non-governmental organizations at any time during the crisis. This data enables us to estimate both the timing and magnitude of the economic shocks experienced by respondents' households between waves. We measure direct income shocks in two distinct ways, and we demonstrate that they yield comparable results. In our main specification, we consider whether respondents have lost more than 20% of their income (combining income from work

²⁷In addition to asking in most waves whether respondents incurred any economic or health shock, in the last wave, we asked them to report the exact amount of household income for every month as well as if they knew anyone hospitalized each month. This allows us to have a more granular and quantifiable measure of economic shock beyond the timing of our survey waves

and other sources) between any two months from February (the pre-shock baseline level) to October 2020, to capture the effects of a sudden large drop in income.

$$shock_1 = \begin{cases} 1, & \text{if } \frac{income_t - income_{t-1}}{income_{t-1}} \leq -0.20 \\ 0, & \text{otherwise} \end{cases}$$

. About 38% of respondents lost at least 20% of their household income, between any two months, between February and October 2020²⁸.

In addition to measures of personal economic shocks, we also control for indirect economic shocks since it is possible that many Americans changed their preferences just by mere exposure to the crisis, such as by knowing someone who was affected economically by the crisis or living in an area that suffered relatively higher economic distress compared to others (Dyer, 2020; Wright et al., 2020). We use data collected and updated in real-time by Harvard’s Opportunity Insights team on the weekly percentage variations in consumer expenditures with respect to the first week of January 2020 (Chetty et al., 2020). This variable is seasonally adjusted and is available at the county level, which we match with the respondents’ residential information²⁹.

Health shocks. Our main measure of direct health shock is whether respondents had a family member, a friend, or an acquaintance who was hospitalized with COVID-19³⁰. About 30% of our respondents knew someone (among family, friends, or acquaintances) who was hospitalized with COVID-19, while 69% knew someone who tested positive. About 33% tested positive for COVID-19 themselves.

We complement these survey-based measures with data on COVID-19 cases in the respondents’ county of residence. While this measure might be subject to different data collection and testing regimes across States, these figures were likely to be the same ones reported by the local media. To improve accuracy, we consider COVID-19 cases³¹ at the county level reported by the middle of each week. We then consider the population size at the county level in 2019 and construct the following measure of increase in cases between week t and $t-1$ in county c : $\frac{cases_{ct} - cases_{ct-1}}{population_{c2019}}$ ³². When, instead, we consider an outcome that is not in

²⁸Respondents who lost at least 20% of their household income between any two months from March to October 2020 are more likely to be young, with a low baseline income, and to belong to a racial minority group. Furthermore, women, Democrats, and those who live in a metropolitan area have incurred such a negative income shock with a significantly higher probability, while cohabitation (or marriage) seems to smooth the financial impact of the pandemic. We control for all these characteristics in our analysis and show how using different specifications does not change our main results

²⁹The Opportunity Insights team uses anonymized data from several private companies to construct public indices of consumer spending, employment, and other outcomes. See Chetty et al. (2020) for further information on series construction.

³⁰We consider this combined measure, as 2.4% of the respondents have a family member who has been hospitalized, 9.8% have a relative, 14.1% a friend, and 14.9% an acquaintance. To control for additional direct health shocks, we also asked respondents their type of health insurance (e.g., public or private), whether they have caring responsibilities towards an elderly or someone with disabilities, who are at greater risk of complications from contracting the virus, and if they knew a healthcare professional who had been working closely with COVID-19 patients

³¹We exploited the data collected by the New York Times from state and local governments and health departments, and available here <https://github.com/nytimes/covid-19-data>.

³²We multiply this measure by 100 to ease the interpretation of the coefficients in our regressions

change, we focus on the logarithm of the cumulative number of cases weighted by the county population: $\ln \left(\frac{cases_{ct}}{population_{c2019}} * 100,000 \right)$.

Media consumption We consider all the news sources consumed by respondents and assess whether they are overall slanted in a direction or not. Therefore, we collected information on respondents’ preferred news sources (including international news and social media) and the number of hours of news they consumed³³. Based on the news sources they indicated, we constructed a “slant score” using the “*AllSides.com*” platform, one of the most commonly used sources of partisan media slant analysis³⁴. The website assigns a score from 1 (Extremely Left) to 5 (Extremely Right) to major sources of news by analyzing their written content on a regular basis and running polls. Matching the scores by Allsides³⁵ to the respondents’ choices, we create an index summing the scores of each source consulted by an individual and divided by the maximum number of possible points.

$$Media\ slant\ index\ (N\ sources\ of\ news) = \frac{\sum_{n=1}^N score_n}{N}$$

This variable measures how politically homogeneous the news sources that respondents consumed are, by taking any value between 1 and 5: the closer a respondent is to 1, the more they consume homogeneous (i.e., less politically diversified) left-leaning media, while the closer they are to 5 the more homogeneous and right-leaning is their media consumption. A score towards the middle indicates either that respondents consume unbiased news, or that they consume news that is slanted in both directions and so that they are exposed to both partisan narratives. Based on this specification, we see that 51% of the Republicans consume Republican-leaning news, and 46% of the Democrats consume Democratic-leaning news. Among independents and non-voters, around 25% (24%) consume Republican-leaning news (Democratic-leaning news).

Other controls. The baseline controls for the regressions are: Gender, Age (in four generations: Gen Z, Gen X, Millennials, Baby Boomers, and above), Education (in four groups: Less than high school, High school or equivalent, Some college, Bachelor or more), Income (in quintiles), Race (White, African Americans, Hispanics, and other), whether respondents are cohabitating with someone else, whether they have children under the age of 18, whether they were in financial difficulties already prior the pandemic, whether they are in the labor force, whether they were unemployed in February (before the pandemic), the U.S. region (South, North-East, Mid-west, West), whether they live in a metropolitan area, whether they have health insurance and whether it’s public or private, the population density of their zip code of residence, whether they are Catholic, or Protestant, and whether they have caring responsibilities. We also control for whether respondents completed the survey in a shorter time than the 99th percentile.


³³The question asked: “Do you get your news from any of these sources (either on television or on the internet)?”, and the multiple option answers were: “ABC, CBS, NBC, CNN, Fox News, MSNBC, and ‘other, please specify’” (e.g., some respondents added The NY Times, The Washington Post, BBC, NPR, and PBS). While there is no exact methodology to measure the partisan slant of news sources (Budak et al., 2016; Groseclose and Milyo, 2005), and since within each source different programs might cover the same news in different tones (Bursztyn et al., 2020), we measured whether respondents were exposed to different points of view during the crisis.

³⁴<https://www.allsides.com/unbiased-balanced-news>

³⁵We use the scores of the first week of April 2020, our baseline wave.


A4.0.3 Public institution media coverage

Figure A2: Media coverage about the CDC



A4.0.4 Survey timing

Figure A3: Timing of the longitudinal survey waves against health indicators of the COVID-19 pandemic



Notes: The figure shows the timing of the seven survey waves implemented between early April and the end of October 2020 against the curves of the 7-day rate of COVID-19 cases and deaths per 100,000 inhabitants, allowing for comparisons between areas with different population sizes in the United States. The figure is based on the U.S. Centers for Disease Control and Prevention publicly available data.

A4.0.5 Sample characteristics

Table A8: Summary statistics of the key demographics in each survey wave, applying survey weights.

	Wave1	Wave2	Wave3	Wave4	Wave5	Wave6	Wave7
Republican	27.37	25.33	26.24	25.94	25.55	24.53	26.47
Democrat	37.93	37.19	35.87	37.41	36.53	36.2	35.94
Independent & other	34.71	37.48	37.88	36.64	37.92	39.27	37.59
Woman	51.69	51.68	51.68	51.69	51.68	51.68	51.7
Age: 18-29	20.51	20.5	20.5	20.5	20.5	20.5	20.53
Age: 30-44	24.71	24.83	25.52	25.97	25.36	25.45	25.5
Age: 45-59	25.02	24.9	24.21	23.76	24.38	24.29	24.22
Age: 60+	29.76	29.77	29.77	29.77	29.77	29.77	29.75
Less than HS	9.77	9.77	9.77	9.77	9.77	9.77	9.77
High school	28.26	28.26	28.26	28.26	28.26	28.26	28.25
Some college	27.71	27.7	27.7	27.7	27.69	27.7	27.73
Bachelor +	34.27	34.27	34.27	34.27	34.28	34.27	34.26
I income q	22.63	22.89	23.89	24.37	23.5	24.41	26.35
II income q	21.35	21.15	20.47	20.4	21.48	19.99	17.86
III income q	17.99	17.28	17.23	17.65	16.39	17.64	17.9
IV income q	18.98	19.37	19.6	18.66	19.7	19.29	18.69
V income q	19.05	19.3	18.81	18.93	18.93	18.66	19.21
Financial hardship pre-COVID-19	31.52	31.04	31.89	31.33	31.14	31.75	31.38
African American	11.93	11.93	11.93	11.93	11.93	11.93	11.93
Hispanic	16.67	16.67	16.67	16.67	16.67	16.67	16.66
Other Race	8.59	8.58	8.58	8.58	8.58	8.58	8.62
White	62.81	62.82	62.82	62.82	62.82	62.82	62.79
Cohabiting	54.58	57.9	57.36	57.98	58.55	62.43	62.68
Parent of minor	27.87	26.64	27.57	27.13	26.68	26.61	27.07
Caring responsibilities	16.48	15.95	16.65	16.21	15.87	16.42	16.37
Not in the labor force	27.68	27.14	28.18	27.75	26.91	27.24	26.99
Unemployed in Feb	9.81	8.96	10.8	9.73	9.06	8.83	9.56
North-East	17.45	17.45	17.45	17.45	17.45	17.45	17.44
Midwest	20.74	20.74	20.74	20.74	20.74	20.74	20.73
South	37.98	37.97	37.97	37.97	37.97	37.97	38
West	23.84	23.84	23.84	23.84	23.84	23.84	23.83
Metropolitan area	85.49	85.1	86.07	84.78	85.29	86.07	85.49
No health insurance	6.67	6.39	6.54	5.89	6.6	6.29	6.24
Population density (ZCTA)	371350.9	368034.6	379836.2	380862.6	380720.3	373379.7	383178.7
<i>N</i>	1441	1228	1177	1219	1199	1192	1076

Table A9: Attrition predictors.

VARIABLES	(1) Attrition W1-W7	(2) Attrition W6-W7		cont.(1) Attrition W1-W7	cont. (2) Attrition W6-W7
Treated		-0.00375 (0.0214)	Cohabiting	-0.309*** (0.0256)	-0.0847*** (0.0272)
$\frac{\log(Covid\ cases\ w7)}{100,000}$	-0.0347* (0.0209)	-0.0424** (0.0212)	Parent	0.0818*** (0.0280)	-0.0190 (0.0295)
Woman	0.0371* (0.0216)	0.0404* (0.0218)	Caring resp.	0.00203 (0.0311)	0.0398 (0.0323)
Age: 18-29	0.0601 (0.0390)	0.107*** (0.0411)	Not labor force	0.0567* (0.0301)	0.0498 (0.0307)
Age: 30-44	-0.00882 (0.0367)	0.0239 (0.0378)	Unemployed in Feb	0.0888* (0.0483)	0.0656 (0.0521)
Age: 45-59	0.0191 (0.0353)	0.0341 (0.0358)	Midwest	-0.0403 (0.0362)	-0.0251 (0.0362)
High school	-0.127* (0.0735)	-0.0940 (0.0774)	South	-0.00920 (0.0349)	-0.0224 (0.0357)
Some college	-0.0885 (0.0714)	-0.0346 (0.0757)	West	0.0135 (0.0364)	-0.000535 (0.0371)
Bachelor +	-0.171** (0.0723)	-0.0929 (0.0765)	Metro area	-0.00353 (0.0327)	0.0182 (0.0322)
II income q	0.219*** (0.0352)	0.147*** (0.0353)	Gov health insur	0.0201 (0.0285)	-0.00581 (0.0293)
III income q	0.283*** (0.0353)	0.185*** (0.0370)	No health insur	0.0756* (0.0454)	0.0893* (0.0488)
IV income q	0.300*** (0.0374)	0.188*** (0.0372)	Pop density	-7.32e-07 (1.06e-06)	-6.11e-07 (9.94e-07)
V income q	0.307*** (0.0380)	0.145*** (0.0371)	Catholic	0.00170 (0.0294)	0.000326 (0.0307)
Financial issues	0.0136 (0.0256)	-0.00475 (0.0262)	Protestant	0.0487* (0.0266)	0.0393 (0.0266)
African American	0.0589 (0.0406)	0.0699* (0.0424)	Rep news	-0.00566 (0.0290)	0.00735 (0.0296)
Hispanic	0.108*** (0.0358)	0.105*** (0.0373)	Dem news	-0.0290 (0.0269)	-0.0286 (0.0264)
Other Race	0.00834 (0.0363)	0.0137 (0.0374)	Constant	0.480*** (0.179)	0.374** (0.184)
Observations	1,442	1,193		1,442	1,193
R-squared	0.185	0.088		0.185	0.088
Mean dep. var.	0.254	0.170		0.254	0.170

Robust standard errors in parentheses; *** p<0.01, ** p<0.05, * p<0.1

A4.0.6 Motivating evidence

Fact I: Institutional trust has declined and has not recovered




Figure A4: Variation in self-reported institutional trust (on a 5-point likert scale, with 1 being the lowest and 5 the highest score) across waves.

Fact II: Trust in public health institutions predicts the adoption of preventive behaviors both in levels and trends, while trust in the government does not. List of preventive behaviors:

- I don't leave home unless strictly needed
- I sneeze or cough into a tissue or the inside of my elbow
- I use pick-up or delivery options more frequently
- I keep a safe distance from others
- I try not to touch objects when I am outside
- I wash my hands more thoroughly and more frequently
- I cover my mouth in public (with a mask, a scarf, ...)

Table A10: Adoption of preventive behaviors and institutional trust

	(1) % prev. behav. adopted	(2) % prev. behav. adopted	(3) % prev. behav. adopted	(4) Δ % prev. behav. adopted w1 -w3	(5) Δ % prev. behav. adopted w1-w3	(6) Δ % prev. behav. adopted w1-w3	(7) Use mask w7	(8) Use mask w7
Trust index PIs w1	0.198*** (0.0728)		0.173** (0.0756)					
Trust index gov w1		0.114** (0.0532)	0.0638 (0.0531)					
Δ trust index PIs w1-w4				0.144*** (0.0500)		0.117** (0.0582)		
Δ trust index gov. w1-w4					0.0661 (0.0790)	0.0288 (0.0878)		
Trust index PIs w7							0.266** (0.129)	0.344** (0.134)
Trust index gov w7								-0.264*** (0.0909)
Constant	0.508*** (0.0599)	0.570*** (0.0530)	0.491*** (0.0632)	0.129* (0.0679)	0.150** (0.0657)	0.150** (0.0680)	0.549*** (0.147)	0.626*** (0.142)
Observations	1,421	1,433	1,417	1,082	1,088	1,073	1,063	1,059
R-squared	0.175	0.168	0.176	0.112	0.104	0.108	0.216	0.230
Demographics	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Covid shock	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
News	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Average dep. var.	0.709	0.709	0.709	0.0344	0.0344	0.0344	0.872	0.872

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In col (1)-(3) the outcome is the fraction of the preventive behaviors adopted from the list above; in cols (4-6), the outcomes is the variation in this fraction between the first and the third week of April; in cols (7) and (8), we flag whether respondents declare to wear a mask in May. All regressions take into account population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Fact III: Perceived politicization predicts drops in trust, both self-reported and factual

Table A11: Perceived polarization and trust

	(1)	(2)	(3)	(4)	(5)	(6)	(7)
	Decrease trust sci. comm. (w7-w1)	Self-reported decrease trust sci. comm.	Decrease trust hospitals (w1-w7)	Self-reported decrease trust hospitals	Decrease trust health insur. (w7-w1)	Self-reported decrease trust health insur. (w7-w1)	Self-reported decr. trust CDC
Sci. comm. politicized	0.122** (0.0503)	0.372*** (0.0370)					
Hospitals politicized			0.160*** (0.0400)	0.264*** (0.0448)			
Health insur. politicized					0.0866** (0.0397)	0.180*** (0.0398)	
CDC politicized							0.399*** (0.0338)
Rep. news	0.0857 (0.0540)	0.132*** (0.0305)	0.129*** (0.0473)	0.0583** (0.0295)	0.0625 (0.0601)	-0.0287 (0.0363)	0.104*** (0.0388)
Dem. news	-0.0275 (0.0490)	-0.0185 (0.0281)	0.0168 (0.0477)	-0.0239 (0.0234)	0.0527 (0.0434)	-0.0117 (0.0409)	0.0849** (0.0351)
Constant	0.258 (0.172)	0.113 (0.103)	0.337** (0.162)	0.133 (0.115)	0.0745 (0.159)	-0.0203 (0.138)	0.0853 (0.120)
Observations	1,004	1,012	1,009	1,014	1,008	1,009	1,005
R-squared	0.091	0.243	0.117	0.166	0.132	0.114	0.249
Demographics	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Covid shocks	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Average decrease	0.286	0.148	0.306	0.0827	0.284	0.187	0.287

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In col. (1) the outcome is a dummy = 1 if the respondent decreased their confidence in the scientific community between wave 7 and wave 1; in col. (2) the outcome is a dummy flagging the self-reported variation in confidence in the scientific community: the dummy =1 if the respondent reports having decreased their confidence in the scientific community from the beginning of the pandemic. Cols (3)-(6) repeat the same exercise but focus on hospitals and healthcare professionals (cols. 3-4) and health insurance companies (cols. 5-6). The outcome of col. (7) is a dummy = 1 if the respondent declared having decreased their confidence in the CDC. All regressions take into account population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Fact IV: Media consumption predicts the perception of the Covid-19 crisis

In the fifth wave of the survey (week of May 18th), we asked respondents to forecast the COVID-19 death rate in the U.S. by the end of the year, after presenting them with the latest official death rate from the CDC and asked for their judgment on how the government handled the crisis³⁶. Our goal was to understand the impact of partisanship on perceptions, partially controlling for a possible information gap by providing the latest official statistics³⁷.

³⁶The questions asked: *By May 17, the U.S. Centers for Disease Control and Prevention (CDC) stated that about 90,000 Americans have so far died from COVID-19 (coronavirus). In addition to this, how many more Americans do you think will die by the end of this year due to coronavirus?*

³⁷Gaines et al. (2007) studies a similar setting showing results of a survey where Americans were asked to state the need and support for the Iraqi war in 2003: while the majority of all respondents thought it was unlikely that the U.S. would ever find weapons of mass destruction, Democrats were more likely to conclude that they simply did not exist while Republicans were more likely to state that they believed the Iraqi government moved or destroyed the weapons.




Figure A5: Expected additional death by news slant: Binscatter, controlling for a set of controls.

Notes: The figure shows a binned scatterplot in which the x-axis variable (news slant) is grouped into equal-sized bins, and the means of the x- and y-axis within each bin are computed. The plot controls for the set of variables described in the notes of table A12.

Table A12: News consumption predicts different perceptions of the pandemic

	(1) Expected additional deaths	(2) Expected additional deaths	(3) Expected additional deaths
Index news slant	-51,825*** (5,037)		
Democratic leaning news		14,624*** (2,790)	14,393*** (2,799)
Republican leaning news		-15,305*** (2,763)	-14,930*** (2,771)
Any direct shock Apr-May	4,772* (2,860)	5,727** (2,825)	5,168* (2,782)
Few COVID-19 news			-7,030*** (2,056)
Constant	54,109*** (9,912)	30,236*** (10,110)	36,097*** (10,502)
Observations	1,184	1,184	1,184
R-squared	0.195	0.210	0.218
Controls	Yes	Yes	Yes
Mean dep. var.	49827	49827	49827

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In col (1)-(3) the outcome is the expected number of additional deaths from COVID-19 that the respondents were expecting would occur in the U.S., between May and the end of the year. All regressions take into account population survey weights and the sampling procedure. The control variables are specified in the Appendix.

A4.0.7 Balance Table




Figure A6: The histogram shows the randomization (null) distribution of the F-statistic obtained by repeatedly reassigning the treatment status according to the experimental design and re-estimating a logit of treatment assignment on baseline covariates, following Kerwin et al. (2024). The red dashed line indicates the F-statistic computed from the observed assignment. This test evaluates whether baseline covariates jointly predict treatment assignment more strongly than would be expected by chance under randomization. Because the observed F-statistic lies well within the bulk of the randomization distribution, we fail to reject the null hypothesis of joint orthogonality between treatment assignment and baseline covariates, indicating no evidence of imbalance beyond chance variation.

Balance table

Table A13: Balance table.

Variable	(1) Control group	(2) Treatment group	(3) Difference
Woman	0.483 (0.500)	0.479 (0.500)	-0.004 (0.029)
Age: 18-29	0.207 (0.406)	0.216 (0.412)	0.008 (0.024)
Age: 30-44	0.282 (0.450)	0.286 (0.452)	0.004 (0.026)
Age: 45-59	0.219 (0.414)	0.224 (0.417)	0.006 (0.024)
High school	0.150 (0.357)	0.159 (0.366)	0.009 (0.021)
Some college	0.392 (0.488)	0.400 (0.490)	0.008 (0.028)
Bachelor +	0.424 (0.495)	0.414 (0.493)	-0.010 (0.029)
II income q	0.188 (0.391)	0.190 (0.392)	0.002 (0.023)
III income q	0.189 (0.392)	0.209 (0.407)	0.019 (0.023)
IV income q	0.194 (0.396)	0.219 (0.414)	0.025 (0.023)
V income q	0.215 (0.411)	0.191 (0.394)	-0.024 (0.023)
Financial hardship	0.287 (0.453)	0.255 (0.436)	-0.032 (0.026)
African American	0.100 (0.300)	0.095 (0.293)	-0.005 (0.017)
Hispanic	0.150 (0.357)	0.152 (0.359)	0.002 (0.021)
Other Race	0.116 (0.320)	0.128 (0.334)	0.012 (0.019)
Cohabiting	0.626 (0.484)	0.614 (0.487)	-0.013 (0.028)
Parent of minor	0.263 (0.440)	0.286 (0.452)	0.024 (0.026)
Caring responsibilities	0.152 (0.359)	0.169 (0.375)	0.017 (0.021)
Not in the labor force	0.256 (0.437)	0.248 (0.432)	-0.008 (0.025)
Unemployed in Feb	0.059 (0.235)	0.053 (0.225)	-0.005 (0.013)
Midwest	0.268 (0.443)	0.257 (0.437)	-0.011 (0.026)
South	0.357 (0.480)	0.352 (0.478)	-0.006 (0.028)
West	0.235 (0.424)	0.234 (0.424)	-0.000 (0.025)
Metropolitan area	0.856 (0.351)	0.876 (0.330)	0.019 (0.020)
Government provided health	0.359 (0.480)	0.338 (0.473)	-0.021 (0.028)
No health insurance	0.073 (0.261)	0.078 (0.268)	0.004 (0.015)
Population density in ZCTA	3,895.589 (9,688.460)	3,765.419 (8,374.398)	-130.169 (525.591)
Catholic	0.171 (0.377)	0.191 (0.394)	0.020 (0.022)
Protestant	0.276 (0.447)	0.278 (0.448)	0.002 (0.026)
Any direct shock Apr-Jun	0.276 (0.447)	0.295 (0.456)	0.019 (0.026)
log county cases/100,000	6.117 (0.954)	6.129 (0.957)	0.012 (0.055)
Var consumer spending by	-0.097 (0.104)	-0.106 (0.120)	-0.009 (0.007)
Observations	613	580	1,442

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1. OLS regression coefficients of the balance table for the survey experiment.

A4.0.8 Additional Results

Table A14: Seconds spent on treatment questions (number of deaths in the state and in the U.S.).

VARIABLES	(1) Seconds spent on treatment questions	(2) Seconds spent on treatment questions	(3) Seconds spent on treatment questions
Informative T	13.97** (5.375)	14.57** (5.766)	15.06*** (4.581)
Constant	54.05*** (4.185)	90.61*** (33.20)	82.09* (46.39)
Observations	1,167	1,167	1,167
R-squared	0.017	0.095	0.145
Demographics	No	Yes	Yes
Covid shock	No	Yes	Yes
News	No	Yes	Yes
Controls x Treat	No	No	Yes
Mean dep. var.	58.95	58.95	58.95

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. The outcome is the time spent answering the questions with the treatment intervention. The last column presented the Lin-adjustment, with the interaction between all the de-meaned controls and the treatment dummy. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Table A15: Two-Stage Least Squares instrumenting correctly selecting the number of COVID-19 deaths in the US and at the State level with the treatment status. Effects on trust and long-term knowledge.

	(1)	(2)	(3)	(4)	(5)	(6)
	TSLS	TSLS	TSLS	TSLS	TSLS	TSLS
VARIABLES	Correct US death rank	Correct US death rank	Incr. trust in CDC	Incr. trust in CDC	CDC not infl. by politics	CDC not infl. by politics
Correct deaths	0.493 (0.389)	0.464* (0.276)	0.655 (0.445)	0.510* (0.295)	0.745 (0.474)	0.551* (0.280)
Constant	0.374** (0.144)	0.571 (0.544)	-0.0804 (0.157)	-0.115 (0.432)	-0.164 (0.166)	-0.657* (0.367)
Observations	953	953	949	949	960	960
Demographics	No	Yes	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
First-stage F	6.245	13.50	6.245	13.50	6.245	13.50
First-stage p	0.0135	0.0003	0.0135	0.0003	0.0135	0.0003
Mean dep. var.	0.550	0.550	0.162	0.162	0.113	0.113

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. We instrument correctly answering the knowledge accuracy question in June with the treatment status. In cols. (1-2), the outcome is whether respondents correctly ranked the U.S. in terms of COVID-19 deaths; in cols (3-4), whether respondents report having increased their confidence in the CDC since the beginning of the pandemic; in cols (5-6), whether respondents report that the CDC is not influenced by politics. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Table A16: Treated respondents are more likely to have tested for Covid-19 and have adopted at least a preventive behavior, but no significant effects on mask-wearing.

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
	Tested for Covid-19	Tested for Covid-19	Adopts no prev. behav.	Adopts no prev. behav.	%Prev. behav. adopted	%Prev. behav. adopted	Covers mouth in public	Covers mouth in public
Informative T	0.0798** (0.0366)	0.0659** (0.0309)	-0.0241** (0.0108)	-0.0234** (0.00933)	0.00236 (0.0222)	-0.000968 (0.0197)	0.0215 (0.0289)	0.0194 (0.0231)
Constant	0.281*** (0.0255)	-0.0493 (0.277)	0.0242*** (0.00886)	-0.0625 (0.0977)	0.750*** (0.0140)	0.861*** (0.143)	0.882*** (0.0228)	0.824*** (0.167)
Observations	987	987	986	986	990	990	986	986
R-squared	0.008	0.099	0.045	0.084	0.064	0.152	0.053	0.188
Demographics	No	Yes	No	Yes	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.331	0.331	0.0188	0.0188	0.735	0.735	0.872	0.872

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In cols (1-2) the outcome is whether respondents declared to have tested for Covid-19; in cols (3-4) whether they declared to have adopted no preventive behavior between Jun and October; in cols (5-6) the fraction of preventive behaviors adopted; in cols (7-8), whether respondents declared to have covered their mouth when in public in June-October. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Table A17: Two-Stage Least Squares instrumenting correctly selecting the number of COVID-19 deaths in the US and at the State level with the treatment status. Effects on preventive behaviors.

VARIABLES	(1) TSLS Tested for Covid-19	(2) TSLS Tested for Covid-19	(3) TSLS Adopts no Prev. Behav.	(4) TSLS Adopts no Prev. Behav.	(5) TSLS % Adopted Prev. Behav.	(6) TSLS % Adopted Prev. Behav.	(7) TSLS Covers mouth In Public	(8) TSLS Covers mouth In Public
Correct deaths	0.722 (0.518)	0.458* (0.272)	-0.196 (0.135)	-0.152* (0.0845)	-0.110 (0.237)	-0.0421 (0.141)	0.112 (0.293)	0.183 (0.181)
Constant	0.0610 (0.180)	-0.404 (0.412)	0.0813 (0.0510)	0.0838 (0.139)	0.794*** (0.0848)	0.841*** (0.218)	0.857*** (0.111)	0.503* (0.278)
Observations	958	958	958	958	960	960	958	958
R-squared						0.181	0.085	0.207
Demographics	No	Yes	No	Yes	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.331	0.331	0.0188	0.0188	0.735	0.735	0.872	0.872
First-stage F	6.245	13.50	6.245	13.50	6.245	13.50	6.245	13.50
First-stage p	0.0135	0.000332	0.0135	0.000332	0.0135	0.000332	0.0135	0.000332

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

TSLS regression coefficients instrumenting knowledge accuracy with the treatment status. In cols (1-2) the outcome is whether respondents declared to have tested for Covid-19; in cols (3-4) whether they declared to have adopted no preventive behavior between Jun and October; in cols (5-6) the fraction of preventive behaviors adopted; in cols (7-8), whether respondents declared to have covered their mouth when in public in June-October. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

A4.0.9 Robustness checks

Testing for Covid-19 Vs contracting Covid-19

Table A18: Treated respondents are more likely to test for Covid-19, but not contract the virus.

	(1)	(2)	(3)
	Respondent tested for COVID	Respondent had COVID	Respondent had COVID
Treated	0.0667** (0.0308)	0.0132 (0.0115)	0.0101 (0.00950)
Constant	-0.0498 (0.274)	0.0169** (0.00705)	-0.195*** (0.0723)
Observations	987	990	990
R-squared	0.099	0.003	0.080
Demographics	Yes	No	Yes
Covid shock	Yes	No	Yes
News	Yes	No	Yes
Mean dep. var.	0.320	0.0223	0.0223

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In col (1), the outcome is whether respondents declared to have been tested for Covid-19 between June and October; in cols (2-3), the outcome is whether respondents declared to have contracted Covid-19 between June and October. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Placebo institutions

Table A19: The treatment has no effect on trust in placebo institutions.

	(1)	(2)	(3)	(4)	(5)	(6)
	Incr. trust in scientific comm.	Incr. trust in healthcare inst. (index)	Incr. trust in financial inst. (index)	Scientific comm. not infl. by politics	Healthcare inst. not infl. by politics (index)	Financial inst. not infl. by politics (index)
Treated	0.0374 (0.0352)	0.0307 (0.0216)	-0.00108 (0.0117)	0.0429 (0.0268)	0.0160 (0.0152)	0.00658 (0.00790)
Constant	0.346 (0.266)	0.257 (0.192)	0.147* (0.0887)	-0.0144 (0.196)	0.251** (0.123)	0.0716 (0.0597)
Observations	986	988	988	990	990	990
R-squared	0.173	0.148	0.127	0.105	0.101	0.142
Demographics	Yes	Yes	Yes	Yes	Yes	Yes
Covid shock	Yes	Yes	Yes	Yes	Yes	Yes
News	Yes	Yes	Yes	Yes	Yes	Yes
Mean dep. var.	0.205	0.168	0.0430	0.137	0.0708	0.0184

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In cols. (1-3), the outcome is whether respondents, between April and October, have increased trust in: (1) the scientific community ; (2) healthcare institutions (hospitals and medical staff, and health insurance companies), measured with an index; (3) economic institutions (banks and financial institutions and the private sector), measured with an index. Cols. (4-6) consider the same institutions, but focus on whether they are perceived as independent from politics. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Attrition

Table A20: Heterogeneous effects by attrition predictors on the probability of correctly ranking the U.S. in terms of per-capita Covid-19 deaths.

VARIABLES	(1) Correct US death rank	(2) Correct US death rank	(3) Correct US death rank	(4) Correct US death rank
Treated	0.0479 (0.0380)	0.120** (0.0463)	0.0765** (0.0329)	0.0891** (0.0389)
Low income	0.0843 (0.0694)			
Low income*Treated	0.0245 (0.0965)			
Woman		0.0624 (0.0558)		
Woman*Treated		-0.129* (0.0716)		
Hispanic			0.115 (0.0822)	
Hispanic*Treated			-0.150 (0.0917)	
Age 18-29				0.0593 (0.109)
Age 18-29*Treated				-0.178 (0.110)
Constant	0.941** (0.439)	1.003** (0.433)	0.990** (0.433)	0.985** (0.441)
Observations	977	977	977	977
R-squared	0.082	0.086	0.084	0.086
Demographics	Yes	Yes	Yes	Yes
Covid shock	Yes	Yes	Yes	Yes
News	Yes	Yes	Yes	Yes

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. The outcome is whether respondents correctly ranked the U.S. in terms of Covid-19 deaths. We consider whether the top 4 predictors of attrition display heterogeneous treatment effects that might drive our results. All regressions take into account population survey weights, the sampling procedure and attrition. The control variables are reported in the Appendix.

Table A21: Heterogeneous effects by attrition predictors on the probability that respondents declare to have increased their trust in the CDC.

VARIABLES	(1) Incr. trust in CDC	(2) Incr. trust in CDC	(3) Incr. trust in CDC	(4) Incr. trust in CDC
Treated	0.0654** (0.0329)	0.0642* (0.0373)	0.0491 (0.0323)	0.0528 (0.0344)
Low income	-0.0430 (0.0573)			
Low income*Treated	-0.0143 (0.0541)			
Woman		0.00242 (0.0309)		
Woman*Treated		-0.00478 (0.0445)		
Hispanic			-0.0525 (0.0486)	
Hispanic*Treated			0.0865 (0.0837)	
Age 18-29				0.00938 (0.0571)
Age 18-29*Treated				0.0458 (0.0675)
Constant	0.438 (0.276)	0.390 (0.277)	0.419 (0.279)	0.404 (0.283)
Observations	978	978	978	978
R-squared	0.088	0.088	0.090	0.088
Demographics	Yes	Yes	Yes	Yes
Covid shock	Yes	Yes	Yes	Yes
News	Yes	Yes	Yes	Yes

Notes: Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$. OLS regression coefficients. The outcome is whether respondents report having increased their confidence in the CDC since the beginning of the pandemic. We consider whether the top 4 predictors of attrition display heterogeneous treatment effects that might drive our results. All regressions take into account population survey weights, the sampling procedure and attrition. The control variables are reported in the Appendix.

Table A22: Heterogeneous effects by attrition predictors on the probability that respondents declare that the CDC is independent from politics.

VARIABLES	(1) CDC not infl. by politics	(2) CDC not infl. by politics	(3) CDC not infl. by politics	(4) CDC not infl. by politics
Treated	0.0496 (0.0327)	0.0426 (0.0379)	0.0381 (0.0283)	0.0466 (0.0291)
Low income	-0.00818 (0.0525)			
Low income*Treated	0.0418 (0.0682)			
Woman		-0.0168 (0.0288)		
Woman*Treated		0.0341 (0.0456)		
Hispanic			-0.109** (0.0477)	
Hispanic*Treated			0.150** (0.0640)	
Age 18-29				0.00697 (0.0767)
Age 18-29*Treated				0.0690 (0.0804)
Constant	-0.0568 (0.194)	-0.0457 (0.183)	-0.0102 (0.185)	-0.0356 (0.189)
Observations	990	990	990	990
R-squared	0.136	0.136	0.142	0.137
Demographics	Yes	Yes	Yes	Yes
Covid shock	Yes	Yes	Yes	Yes
News	Yes	Yes	Yes	Yes

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. The outcome is whether respondents report that politics do not influence the CDC. We consider whether the top 4 predictors of attrition display heterogeneous treatment effects that might drive our results. All regressions take into account population survey weights, the sampling procedure and attrition. The control variables are reported in the Appendix.

Table A23: Heterogeneous effects by attrition predictors on the probability that respondents declare having tested for Covid-19.

VARIABLES	(1) Tested for Covid-19	(2) Tested for Covid-19	(3) Tested for Covid-19	(4) Tested for Covid-19
Treated	0.0844** (0.0369)	0.0481 (0.0419)	0.0718** (0.0312)	0.0794** (0.0357)
Low income	0.0320 (0.0625)			
Low income*Treated	-0.0724 (0.0775)			
Woman		-0.0307 (0.0552)		
Woman*Treated		0.0347 (0.0725)		
Hispanic			0.109 (0.0774)	
Hispanic*Treated			-0.0392 (0.0968)	
Age 18-29				0.0461 (0.0781)
Age 18-29*Treated				-0.0674 (0.0860)
Constant	-0.0260 (0.275)	-0.00788 (0.277)	-0.0263 (0.280)	-0.0319 (0.278)
Observations	987	987	987	987
R-squared	0.098	0.098	0.097	0.098
Demographics	Yes	Yes	Yes	Yes
Covid shock	Yes	Yes	Yes	Yes
News	Yes	Yes	Yes	Yes

Notes: Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

OLS regression coefficients. The outcome is whether respondents declare having tested for Covid-19 in the past months. We consider whether the top 4 predictors of attrition display heterogeneous treatment effects that might drive our results. All regressions take into account population survey weights, the sampling procedure, and attrition. The control variables are reported in the Appendix.

Controlling for interpersonal trust

Table A24: Controlling for interpersonal trust does not change the results.

VARIABLES	(1) Correct US death rank	(2) Correct US death rank	(3) Incr. trust in CDC	(4) Incr. trust in CDC	(5) CDC not infl. by politics	(6) CDC not infl. by politics
Informative T	0.0393 (0.0331)	0.0530* (0.0303)	0.0617* (0.0316)	0.0596** (0.0291)	0.0544* (0.0297)	0.0596** (0.0269)
Most people can be trusted	-0.0337 (0.0482)	0.0122 (0.0426)	-0.0291 (0.0475)	-0.0239 (0.0388)	-0.0317 (0.0435)	-0.0343 (0.0350)
Can't be too careful with people	-0.0555 (0.0439)	-0.0254 (0.0472)	-0.0833* (0.0422)	-0.0781** (0.0380)	-0.00680 (0.0388)	-0.0238 (0.0305)
Constant	0.565*** (0.0340)	1.048** (0.445)	0.173*** (0.0335)	0.423 (0.283)	0.0915*** (0.0280)	-0.0237 (0.187)
Observations	977	977	978	978	990	990
R-squared	0.004	0.083	0.024	0.095	0.023	0.137
Demographics	No	Yes	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.550	0.550	0.162	0.162	0.113	0.113

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In col (1-2) the outcome is whether respondents correctly estimated the number of Covid-19-related deaths in the U.S. and in their State, in June; in cols (3-4), the outcome is whether respondents (in October 2020) report having increased their confidence in the CDC since the beginning of the pandemic; in cols (5-6), the outcome is whether respondents (in October 2020) report that politics do not influence the CDC. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Alternative specifications

Table A25: Repeating the analyses using a probit model does not affect results.

VARIABLES	June (w6)		October (w7)					
	(1) Correct death est.	(2) Correct death est.	(3) Correct US death rank	(4) Correct US death rank	(5) Incr. trust in CDC	(6) Incr. trust in CDC	(7) CDC not infl. by politics	(8) CDC not infl. by politics
Informative T	0.285*** (0.0891)	0.387*** (0.0835)	0.103 (0.0848)	0.144* (0.0812)	0.266** (0.134)	0.272** (0.121)	0.286* (0.164)	0.394** (0.168)
Constant	-0.556*** (0.0636)	-1.235*** (0.451)	0.0708 (0.0738)	1.481 (1.190)	-1.158*** (0.0853)	0.102 (1.125)	-1.398*** (0.125)	-2.999** (1.165)
Observations	1,147	1,146	977	977	971	964	990	984
Demographics	No	Yes	No	Yes	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.330	0.330	0.550	0.550	0.162	0.162	0.113	0.113

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In col (1-2), the outcome is whether respondents correctly estimated the number of Covid-19-related deaths in the U.S. and in their State, in June; in cols (3-4), the outcome is whether respondents (in October 2020) correctly ranked the U.S. in terms of Covid-19 deaths per capita; in cols (5-6), the outcome is whether respondents (in October 2020) report having increased their confidence in the CDC since the beginning of the pandemic; in cols (7-8), the outcome is whether respondents (in October 2020) report that the CDC is not influenced by politics. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Table A26: Repeating the analyses considering alternative measures of the main outcomes confirms the results.

VARIABLES	(1) Decr. trust in CDC	(2) Decr. trust in CDC	(3) CDC is politicized	(4) CDC is politicized	(5) CDC independence scale (1-4)	(6) CDC independence scale (1-4)
Informative T	-0.0220 (0.0373)	-0.0209 (0.0363)	-0.0621* (0.0341)	-0.0658** (0.0298)	0.185** (0.0897)	0.199*** (0.0735)
Constant	0.305*** (0.0246)	-0.416 (0.270)	0.269*** (0.0263)	0.0238 (0.265)	2.139*** (0.0607)	1.797*** (0.593)
Observations	978	978	983	983	983	983
R-squared	0.005	0.110	0.010	0.106	0.017	0.137
Demographics	No	Yes	Yes	Yes	Yes	Yes
Covid shock	No	Yes	Yes	Yes	Yes	Yes
News	No	Yes	Yes	Yes	Yes	Yes
Mean dep. var.	0.280	0.280	0.240	0.240	2.239	2.239

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In col (1-2), the outcome is whether respondents (in October 2020) report having decreased their confidence in the CDC since the beginning of the pandemic; in cols (3-4), the outcome is whether respondents (in October 2020) report that the CDC is very influenced by politics; in cols (5-6), the perception of whether the CDC is influenced by politics measured with a 4-point likert scale (from "very influenced", to "not influenced at all"). All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Validity of meta-cognition

Table A27: Self-reported variation in confidence in an institution predicts the variation in trust recorded in previous surveys over time.

	(1) Variation in confidence in the scientific community	(2) Variation in confidence in the scientific community	(3) Variation in confidence in hospitals and medical staff	(4) Variation in confidence in hospitals and medical staff	(5) Variation in confidence in health insurance companies	(6) Variation in confidence in health insurance companies
Self-reported var. in trust in sci. comm.	0.198*** (0.0412)	0.153*** (0.0394)				
Self-reported var. in trust in hospitals			0.155*** (0.0438)	0.150*** (0.0437)		
Self-reported var. in trust in health insur.					0.0863 (0.0554)	0.0935* (0.0540)
Constant	-0.505*** (0.0871)	0.0460 (0.476)	-0.478*** (0.0967)	-0.348 (0.404)	-0.279*** (0.107)	-0.0468 (0.513)
Observations	1,064	1,064	1,066	1,066	1,060	1,060
R-squared	0.042	0.093	0.017	0.072	0.006	0.095
Demographics	No	Yes	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
Mean dep. var.	-0.0873	-0.0873	-0.139	-0.139	-0.113	-0.113

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In col (1-2), the outcome is the variation in confidence in the scientific community between the surveys run in April and October; in cols (3-4), the outcome is the variation in confidence in hospitals and medical staff between the surveys run in April and October; in cols (5-6), the outcome is the variation in confidence in health insurance companies between the surveys run in April and October. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

Table A28: Self-reported variation in confidence in an institution predicts the variation in trust recorded in previous surveys over time.

	(1) Variation in confidence in banks and financial inst.	(2) Variation in confidence in banks and financial inst.	(3) Variation in confidence in the private sector	(4) Variation in confidence in the private sector
Self-reported var. in trust in fin. inst.	0.202*** (0.0632)	0.191*** (0.0577)		
Self-reported var. in trust in priv. sect.			0.221*** (0.0541)	0.184*** (0.0528)
Constant	-0.555*** (0.125)	-1.139*** (0.407)	-0.480*** (0.107)	0.209 (0.393)
Observations	1,060	1,060	1,061	1,061
R-squared	0.020	0.104	0.032	0.106
Demographics	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes
News	No	Yes	No	Yes
Mean dep. var.	-0.163	-0.163	-0.0719	-0.0719

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In cols (1-2), the outcome is the variation in confidence in banks and financial institutions between the April and October surveys; in cols (3-4), the outcome is the variation in confidence in the private sector between the April and October surveys. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

A4.0.10 Heterogeneous effects by media consumption

Table A29: Heterogeneous effects by media consumption

VARIABLES	(1) Tot. time question	(2) Correct death est.	(3) Correct US death rank	(4) Incr. trust in CDC	(5) CDC not infl. by politics
Informative T	14.04 (9.083)	0.150*** (0.0325)	0.0218 (0.0446)	0.0722* (0.0377)	0.0868** (0.0375)
Rep news	-20.22* (10.45)	-0.0220 (0.0577)	-0.265*** (0.0691)	-0.114*** (0.0374)	-0.00390 (0.0382)
Dem news	3.482 (11.14)	0.0695 (0.0469)	0.0288 (0.0623)	0.0247 (0.0506)	0.0371 (0.0454)
Treated*Dem leaning news	-11.35 (11.89)	-0.0785 (0.0600)	-0.0747 (0.0776)	-0.0590 (0.0744)	-0.0227 (0.0626)
Treated*Rep leaning news	15.87 (13.85)	-0.0257 (0.0662)	0.257*** (0.0872)	0.0207 (0.0625)	-0.0985 (0.0598)
Constant	91.36*** (32.48)	0.367** (0.148)	1.055*** (0.183)	0.116 (0.151)	0.229* (0.137)
Observations	1,167	1,147	977	978	990
R-squared	0.098	0.160	0.097	0.092	0.136
Demographics	Yes	Yes	Yes	Yes	Yes
Covid shock	Yes	Yes	Yes	Yes	Yes
News	Yes	Yes	Yes	Yes	Yes
p-val Tx*Rep News = Tx*Dem News	0.060	0.508	0.000	0.273	0.241
Mean dep. var.	58.95	0.330	0.550	0.162	0.113

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. The outcomes are: col (1) total time in seconds spent answering the treatment questions; col (2) whether respondents correctly estimated the number of Covid-19-related deaths in the U.S. and in their State, in June; col (3) whether respondents (in October 2020) correctly ranked the U.S. in terms of Covid-19 deaths per capita; col (4) whether respondents (in October 2020) report having increased their confidence in the CDC since the beginning of the pandemic; col (5) whether respondents (in October 2020) report that the CDC is not influenced by politics. The omitted news category is respondents whose news diet is balanced or unslanted, so *Informative T* is the treatment effect for that group. The interaction on perceived independence in col (5) has $p = 0.10$. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix. Produced by CDC/Codes/4.RC.do (news heterogeneity block).

Table A30: News consumption predicts time spent answering the question (attention) and whether the answer is correct in the control group.

VARIABLES	(1) Tot. time question	(2) Tot. time question	(3) Correct death est.	(4) Correct death est.
Rep news	-16.19 (10.10)	-25.06** (11.70)	-0.0490 (0.0690)	-0.0161 (0.0588)
Dem news	3.463 (13.12)	-0.595 (9.907)	0.0800 (0.0533)	0.0651 (0.0448)
Constant	57.22*** (8.236)	102.7** (50.54)	0.283*** (0.0353)	0.538** (0.210)
Observations	1,175	1,175	1,168	1,168
R-squared	0.011	0.121	0.015	0.191
Demographics	No	Yes	No	Yes
Covid shock	No	Yes	No	Yes
News	No	Yes	No	Yes
Mean dep. var.	52.13	52.13	0.280	0.280
p-val Dem News = Rep news	0.0752	0.0227	0.0893	0.230

Notes: Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.


OLS regression coefficients. Only the subsample of the control group is considered; the mean of the dependent variable is the survey-weighted control-group mean. The outcomes are: col (1-2) Total time spent answering the treatment question; col (3-4) whether respondents correctly estimated the number of Covid-19-related deaths in the U.S. and in their State, in June. All regressions account for population survey weights and the sampling procedure. The control variables are reported in the Appendix.

A4.1 Study II

A4.1.1 Controls

The controls closely align with those adopted in the previous experiment. The news index is built exactly as above, and the other controls are the same, except for the information on Covid-19, which is not applicable in this case, and the information on health insurance, which was not available for these respondents. Therefore, the final list of controls is: Gender, Age (in four generations: Gen Z, Gen X, Millennials, Baby Boomers, and above), Education (in four groups: Less than high school, High school or equivalent, Some college, Bachelor or more), Income (in quintiles), Race (White, African Americans, Hispanics, and other), whether respondents are cohabitating with someone else, whether they have children under the age of 18, whether they were unemployed the U.S. region (South, North-East, Mid-west, West), whether they live in a metropolitan area, whether they are Catholic, or Protestant. We also control for whether respondents completed the survey in a shorter or longer time than the 99th percentile.

Figure A7: Randomized treatment A (low search costs, high value)




Notes: The figure shows the landing webpage that respondents would see if they clicked on the link.

Figure A8: Randomized treatment B (low search costs, low value)

Regional Offices

We have ten Regional offices, each of which is responsible for the execution of our programs within several states and territories. To get information about activities in your Regional office, select the appropriate link.



- [Region 1](#) – Boston (serving CT, ME, MA, NH, RI, and VT)
- [Region 2](#) – New York City (serving NJ, NY, Puerto Rico, the U.S. Virgin Islands and 8 federally recognized Indian Nations)
- [Region 3](#) – Philadelphia (serving DE, DC, MD, PA, VA, WV and 7 federally recognized tribes)
- [Region 4](#) – Atlanta (serving AL, FL, GA, KY, MS, NC, SC, and TN)
- [Region 5](#) – Chicago (serving IL, IN, MI, MN, OH, and WI)
- [Region 6](#) – Dallas (serving AR, LA, NM, OK, and TX)
- [Region 7](#) – Kansas City (serving IA, KS, MO, and NE)
- [Region 8](#) – Denver (serving CO, MT, ND, SD, UT, and WY)
- [Region 9](#) – San Francisco (serving AZ, CA, HI, NV, American Samoa, Commonwealth of the Northern Mariana Islands, Federated States of Micronesia, Guam, Marshall Islands, and Republic of Palau)
- [Region 10](#) – Seattle (serving AK, ID, OR, WA and 271 native tribes)

Notes: The figure shows the landing webpage that respondents would see if they clicked on the link.

Figure A9: Media coverage about the EPA

'It would be devastating': inside Trump's plan to destroy the EPA

Trump has made campaign promises to toss crucial environmental regulations - including dismantling the federal body with the most power to tackle the climate crisis



POLITICO

How Trump could reshape EPA on 'Day 1'

By FRANCISCO "A.J." CAMACHO | 07/18/2024 06:16 PM EDT

CNN Climate Solutions Weather

(CNN) — Environmental Protection Agency administrator Michael Regan is bullish on his agency's prospects against the harshest headwinds in its history, as Republicans, anti-climate industry and an ultra-conservative Supreme Court take aim at its core mission.

"I feel extremely confident in the product we put out," Regan told CNN in an interview at a Climate Action Campaign event on Tuesday, just hours after Republican attorneys general and industry groups asked the Supreme Court to temporarily block the EPA from enforcing new rules that aim to curb planet-warming emissions from power plants.

The EPA's new rules will compel some coal and natural gas power plants to either cut or capture 90% of their climate pollution by 2032. The rules are expected to reduce the carbon dioxide emissions from the sector by 75% compared to a peak in 2005, and they are a huge part of President Joe Biden's strategy to quickly cut US planet-warming pollution.

Balance table

Table A31: Balance table.

Variable	(1) Mean T1	(2) Mean T2	(3) Mean T3	(4) T1 vs T2	(5) T1 vs T3	(6) T2 vs T3
Knows EPA	0.731 (0.444)	0.679 (0.467)	0.714 (0.452)	0.052* (0.028)	0.017 (0.027)	0.035 (0.028)
Rep news	0.144 (0.352)	0.183 (0.387)	0.193 (0.395)	-0.039* (0.023)	-0.049** (0.023)	0.010 (0.023)
Dem news	0.190 (0.393)	0.185 (0.389)	0.174 (0.379)	0.005 (0.024)	0.017 (0.024)	-0.012 (0.023)
Woman	0.521 (0.500)	0.555 (0.497)	0.533 (0.499)	-0.034 (0.030)	-0.012 (0.030)	-0.022 (0.030)
African	0.125 (0.331)	0.118 (0.323)	0.116 (0.321)	0.007 (0.020)	0.009 (0.020)	-0.002 (0.019)
American	0.152 (0.359)	0.178 (0.383)	0.161 (0.368)	-0.026 (0.023)	-0.009 (0.022)	-0.017 (0.023)
Hispanic	0.085 (0.279)	0.100 (0.300)	0.072 (0.258)	-0.015 (0.018)	0.013 (0.016)	-0.028* (0.017)
Not white, black or hisp	0.142 (0.350)	0.147 (0.354)	0.163 (0.370)	-0.005 (0.022)	-0.020 (0.022)	0.016 (0.022)
Age: 18-29	0.269 (0.444)	0.312 (0.464)	0.297 (0.457)	-0.043 (0.028)	-0.028 (0.027)	-0.015 (0.028)
Age: 30-44	0.246 (0.431)	0.229 (0.420)	0.242 (0.428)	0.017 (0.026)	0.005 (0.026)	0.013 (0.025)
Age: 45-59	0.196 (0.397)	0.171 (0.376)	0.195 (0.397)	0.026 (0.024)	0.001 (0.024)	0.024 (0.023)
HS	0.381 (0.486)	0.388 (0.488)	0.399 (0.490)	-0.008 (0.030)	-0.018 (0.030)	0.011 (0.029)
College	0.371 (0.484)	0.385 (0.487)	0.343 (0.475)	-0.014 (0.030)	0.028 (0.029)	-0.041 (0.029)
BA or above	0.252 (0.435)	0.232 (0.423)	0.233 (0.423)	0.020 (0.026)	0.019 (0.026)	0.000 (0.025)
\$30k-\$60k	0.213 (0.410)	0.269 (0.444)	0.265 (0.442)	-0.055** (0.026)	-0.051** (0.026)	-0.004 (0.027)
\$60k-\$100k	0.313 (0.464)	0.296 (0.457)	0.284 (0.452)	0.018 (0.028)	0.029 (0.028)	-0.011 (0.027)
>\$100k	0.244 (0.430)	0.269 (0.444)	0.274 (0.446)	-0.024 (0.027)	-0.029 (0.027)	0.005 (0.027)
Midwest	0.348 (0.477)	0.301 (0.459)	0.347 (0.476)	0.047 (0.029)	0.001 (0.029)	0.046 (0.028)
South	0.248 (0.432)	0.305 (0.461)	0.252 (0.435)	-0.057** (0.027)	-0.004 (0.026)	-0.053* (0.027)
West	0.612 (0.488)	0.633 (0.482)	0.628 (0.484)	-0.022 (0.030)	-0.016 (0.030)	-0.005 (0.029)
Employed	0.056 (0.230)	0.071 (0.257)	0.068 (0.252)	-0.015 (0.015)	-0.012 (0.015)	-0.003 (0.015)
Unemployed	0.535 (0.499)	0.501 (0.500)	0.503 (0.500)	0.034 (0.031)	0.032 (0.030)	0.002 (0.030)
Widowed, divorced,	0.321 (0.467)	0.376 (0.485)	0.358 (0.480)	-0.055* (0.029)	-0.037 (0.029)	-0.018 (0.029)
Parent of minor	0.937 (0.244)	0.936 (0.244)	0.905 (0.293)	0.000 (0.015)	0.031* (0.016)	-0.031* (0.016)
HH internet access via	0.387 (0.487)	0.383 (0.487)	0.417 (0.493)	0.004 (0.030)	-0.030 (0.030)	0.034 (0.029)
State has Rep. gov.	0.387 (0.487)	0.383 (0.487)	0.417 (0.493)	0.004 (0.030)	-0.030 (0.030)	0.034 (0.029)
State has Rep. gov.	0.885 (0.320)	0.838 (0.368)	0.846 (0.361)	0.046** (0.021)	0.038* (0.021)	0.008 (0.022)
Metropolitan area	0.196 (0.397)	0.162 (0.368)	0.191 (0.394)	0.035 (0.023)	0.005 (0.024)	0.030 (0.023)
catholic	0.231 (0.422)	0.243 (0.429)	0.270 (0.444)	-0.012 (0.026)	-0.039 (0.026)	0.027 (0.026)
protestant	49.231 (20.451)	49.443 (21.448)		-0.212 (1.282)		
AQI Score	0.475 (0.500)	0.465 (0.499)		0.010 (0.031)		
AQI>50						
Observations	520	551	559	1,071	1,079	1,110

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1. OLS regression coefficients of the balance table for the survey experiment.




Figure A10: The histogram shows the randomization (null) distribution of the likelihood-ratio (LR) statistic obtained by repeatedly reassigning treatment labels according to the experimental design and re-estimating a multinomial logit of treatment assignment on baseline covariates, following Kerwin et al. (2024). The red dashed line indicates the LR statistic computed from the observed assignment. This test evaluates whether baseline covariates jointly predict treatment assignment more strongly than would be expected by chance under randomization. Because the observed LR lies well within the bulk of the randomization distribution, we fail to reject the null hypothesis of joint orthogonality between treatment assignment and baseline covariates, suggesting no imbalance beyond chance.

Table A32: Trust outcomes on their original scales.

	(1) EPA improves lives (1–7)	(2) EPA improves lives (1–7)	(3) EPA indepen- dence (1–7)	(4) EPA indepen- dence (1–7)	(5) Support for EPA (0–100)	(6) Support for EPA (0–100)
T1: Low search costs & high value info	0.139* (0.0832)	0.0554 (0.0786)	0.241*** (0.0864)	0.199** (0.0844)	3.190* (1.717)	1.640 (1.628)
T2: High search costs & high value info	0.0327 (0.0818)	0.00607 (0.0772)	0.162* (0.0851)	0.164** (0.0829)	1.163 (1.696)	0.796 (1.605)
Constant	4.764*** (0.0582)	4.386*** (0.240)	3.243*** (0.0605)	3.847*** (0.258)	59.41*** (1.201)	59.72*** (5.003)
Observations	1,626	1,626	1,613	1,613	1,580	1,580
R-squared	0.018	0.154	0.007	0.087	0.003	0.136
Demographics	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
p-val T1 = T2	0.206	0.533	0.369	0.683	0.236	0.598
Mean dep. var.	4.788	4.788	3.360	3.360	60.61	60.61

Notes: Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Columns (3)–(8) of Table 2 with the trust outcomes on the scales as elicited: perceived benevolence and independence on 7-point Likert scales, and the percent likelihood of recommending the EPA to lead a community study. Excluded group: T3. The p-val row uses robust standard errors, as in the published version of this table. Binarising as specified in the pre-analysis plan (main text) does not change the pattern: the high-value link raises perceived independence on both measures, and the benevolence and support effects are significant without controls and attenuate with them on both.

Table A33: Knowledge accuracy at the pre-registered ± 10 window.

VARIABLES	(1)	(2)
	Knowledge accuracy (± 10)	
T1: Low search costs & high value info	0.141*** (0.0323)	0.132*** (0.0323)
T2: High search costs & high value info	-0.208*** (0.0303)	-0.203*** (0.0303)
Constant	0.489*** (0.0224)	0.423*** (0.0967)
Observations	1,400	1,400
R-squared	0.094	0.131
Demographics	No	Yes
News	No	Yes
p-val T1 = T2	0.000	0.000
Mean dep. var.	0.441	0.441

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

Columns (1)–(2) of Table 2 re-estimated with the ± 10 accuracy window specified in the pre-analysis plan. Arm means: T1 0.618, T2 0.271, T3 0.479. T3 is scored on the office-location question and does not depend on the window, so widening it raises T1 and T2 accuracy while leaving T3 fixed; at ± 10 T1 exceeds T3, at ± 5 the two are indistinguishable. The trust columns of Table 2 do not depend on the window. Produced by TESS/Codes/EPA/11.mechanism_tables_w5.do.

Table A34: Trust and opening the link, by treatment arm: original scales.

VARIABLES	(1) EPA improves lives (1–7)	(2) EPA improves lives (1–7)	(3) EPA independent (1–7)	(4) EPA independent (1–7)	(5) Support for EPA (0–100)	(6) Support for EPA (0–100)
T1: Low search costs & high value info	0.148 (0.142)	0.0738 (0.136)	0.146 (0.151)	0.139 (0.155)	0.411 (2.710)	-0.741 (2.628)
Clicked link	0.195* (0.117)	0.166 (0.110)	-0.0873 (0.119)	-0.0633 (0.118)	2.640 (2.459)	2.218 (2.375)
Clicked $\times$ T1	-0.0466 (0.175)	-0.0438 (0.168)	0.160 (0.185)	0.125 (0.184)	3.748 (3.502)	3.358 (3.345)
Constant	4.658*** (0.0901)	4.016*** (0.275)	3.291*** (0.0912)	3.839*** (0.318)	57.95*** (1.851)	53.90*** (6.070)
T1 vs T3, among those who clicked	0.101 (0.102)	0.0300 (0.0984)	0.306*** (0.107)	0.264*** (0.102)	4.159* (2.221)	2.616 (2.103)
T1 vs T3, among those who did not	0.148 (0.142)	0.0738 (0.136)	0.146 (0.151)	0.139 (0.155)	0.411 (2.710)	-0.741 (2.628)
Observations	1,075	1,075	1,067	1,067	1,049	1,049
R-squared	0.024	0.156	0.010	0.087	0.011	0.143
Demographics	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
Mean dep. var.	4.799	4.799	3.346	3.346	60.75	60.75

Notes: Standard errors in parentheses, robust to heteroskedasticity. *** p<0.01, ** p<0.05, * p<0.1.

Table 3 with the trust outcomes on the scales as elicited. Produced by TESS/Codes/EPA/11.mechanism_tables_w5.do.

A4.1.2 Robustness checks

Table A35: Clicked treatment link

VARIABLES	(1) Clicked link	(2) Clicked link	(3) Clicked link
Knows EPA	0.0435 (0.0362)	0.0784 (0.0488)	0.0446 (0.0361)
Informative treatment	0.120*** (0.0293)	0.174*** (0.0557)	0.0893** (0.0366)
Informative T*Dem news			0.135* (0.0778)
Informative T*Rep news			0.0348 (0.0806)
Informative T*Knows EPA		-0.0750 (0.0651)	
Constant	0.525*** (0.111)	0.501*** (0.114)	0.539*** (0.113)
Observations	1,079	1,079	1,079
R-squared	0.086	0.087	0.089
Demographics	Yes	Yes	Yes
News	No	Yes	Yes
Mean dep. var.	0.601	0.601	0.601

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. The outcome is whether respondents clicked the treatment link provided in the low-cost treatment groups (T1 and T3). Only the subsample of the low-cost treatment groups - which were provided a link - is included. The control variables are reported in the Appendix.

Table A36: Pre-registered secondary outcome: clicking an EPA social-media link at the end of the survey.

VARIABLES	(1)	(2)	(3)	(4)	(5)	(6)
	Clicked any EPA social-media link					
T1: Low search costs & high value info	0.0150 (0.0177)	0.0128 (0.0175)			0.0152 (0.0204)	0.0164 (0.0202)
T2: High search costs & high value info	0.0429** (0.0174)	0.0402** (0.0172)			0.0434** (0.0199)	0.0420** (0.0196)
Clicked treatment link			-0.0100 (0.0198)	-0.0140 (0.0198)		
Correct answer					0.0328* (0.0175)	0.0325* (0.0175)
Constant	0.0762*** (0.0124)	0.223*** (0.0534)	0.101*** (0.0106)	0.235*** (0.0645)	0.0684*** (0.0165)	0.205*** (0.0609)
Observations	1,630	1,630	1,110	1,110	1,400	1,400
R-squared	0.006	0.061	0.003	0.053	0.006	0.066
Sample	All	All	T1, T3	T1, T3	All	All
Demographics	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
p-val T1 = T2	0.116	0.119				
Mean dep. var.	0.093	0.093	0.095	0.095	0.093	0.093

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. The outcome equals one if the respondent clicked the EPA's Facebook, X or Instagram link offered after the last survey question (a pre-registered secondary outcome). Shares by arm: 7.3 percent in T3, 8.9 percent in T1, 11.6 percent in T2. Cols (1)–(2): treatment arms, T3 omitted, all respondents. Cols (3)–(4): the two arms that were offered a link, on whether the respondent opened it; clicking is a post-treatment choice, so these are descriptive associations. Cols (5)–(6): all respondents, on whether the factual question was answered correctly (± 5 window for T1 and T2, the regional-office city for T3), with arm indicators; answering correctly is likewise post-treatment. The control variables are reported in the Appendix. Produced by `TESS/Codes/EPA/15.mht.socialmedia.do`.

Knowing the PI

Table A37: Knowing the EPA is correlated with higher levels of trust

	(1)	(2)	(3)	(4)	(5)	(6)
	EPA improves lives	EPA improves lives	EPA indepen- dence scale	EPA indepen- dence scale	Support for EPA's involvement	Support for EPA's involvement
Knows EPA	0.420*** (0.0684)	0.285*** (0.0734)	-0.147** (0.0732)	-0.176** (0.0784)	10.47*** (1.467)	7.634*** (1.579)
Constant	4.516*** (0.0550)	4.216*** (0.228)	3.481*** (0.0593)	4.069*** (0.252)	53.36*** (1.226)	55.20*** (4.685)
Observations	1,626	1,626	1,613	1,613	1,580	1,580
R-squared	0.035	0.160	0.004	0.085	0.033	0.148
Demographics	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
Mean dep. var.	4.788	4.788	3.360	3.360	60.61	60.61

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In cols (1-2), the outcome reports whether the EPA recommends policies and regulations that improve or worsen the life of the respondents and their loved ones on a 7-point likert scale from "Always worsen" to "Always improve"; in cols (3-4), whether the EPA is politicized versus independent on a 7-point likert scale from "Completely politicized" to "Completely independent"; in cols (5-6), the outcome is percent likelihood that the respondents would recommend the EPA to conduct a study to measure air pollution levels in their community. The control variables are reported in the Appendix.

Table A38: Knowing the EPA and perceived political independence, by news slant

VARIABLES	(1) EPA independence scale (1-7)	(2) EPA independence scale (1-7)
Knows EPA	-0.240*** (0.0862)	-0.216** (0.0902)
Rep news	-0.453*** (0.156)	-0.415*** (0.159)
Dem news	-0.0780 (0.171)	-0.0954 (0.173)
Knows EPA × Rep news	-0.376** (0.188)	-0.420** (0.189)
Knows EPA × Dem news	0.768*** (0.204)	0.750*** (0.206)
Constant	3.577*** (0.0681)	4.080*** (0.249)
Knows EPA, among Rep-leaning news consumers	-0.616*** (0.168)	-0.636*** (0.172)
Knows EPA, among Dem-leaning news consumers	0.527*** (0.185)	0.534*** (0.192)
Observations	1,613	1,613
R-squared	0.081	0.098
Demographics	No	Yes
Mean dep. var.	3.360	3.360

Notes: Standard errors in parentheses, robust to heteroskedasticity. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. The outcome is whether the EPA is politicized versus independent on a 7-point Likert scale from “Completely politicized” to “Completely independent”. The omitted news category is respondents whose news diet is balanced or unslanted, so the coefficient on *Knows EPA* is the association for that group; the rows below the line report the implied association among Republican- and Democratic-leaning news consumers. Raw means of the outcome: 2.49 among Republican-leaning news consumers who know the EPA, 3.10 among those who do not, 3.49 and 3.54 among all other respondents who do and do not know it. The share consuming Republican-leaning news is 16.5 percent among respondents who know the EPA and 19.7 percent among those who do not. The control variables are reported in the Appendix. Produced by TESS/Codes/EPA/14.knows_epa_by_news.do.

Table A39: Predictors of knowing the PI

	(1) Knows EPA
Rep news	-0.0576* (0.0306)
Dem news	0.00276 (0.0266)
Woman	-0.139*** (0.0211)
African American	-0.143*** (0.0374)
Hispanic	-0.118*** (0.0330)
Not white, black or hisp	-0.109*** (0.0418)
Age: 18-29	-0.207*** (0.0426)
Age: 30-44	-0.172*** (0.0319)
Age: 45-59	-0.0990*** (0.0302)
HS	0.156*** (0.0590)
College	0.259*** (0.0567)
BA or above	0.347*** (0.0577)
\$30k-\$60k	-0.0121 (0.0343)
\$60k-\$100k	0.0183 (0.0351)
>\$100k	0.0479 (0.0367)
Midwest	-0.0522 (0.0349)
South	-0.0460 (0.0364)
West	-0.101*** (0.0340)
Employed	0.0115 (0.0266)
Unemployed	-0.0356 (0.0503)
Widowed, divorced, separated, never married	0.0745*** (0.0244)
Parent of minor	-0.00250 (0.0265)
HH internet access	0.0591 (0.0455)
State has Rep. gov.	-0.00998 (0.0251)
Metropolitan area	-0.0284 (0.0308)
Catholic	-0.0241 (0.0295)
Protestant	0.0511** (0.0255)
Constant	0.661*** (0.0828)
Observations	1,604
R-squared	0.157
Demographics	Yes
News	Yes
Mean dep. var.	0.707

Notes: Standard errors in parentheses. ***
p<0.01, ** p<0.05, * p<0.1. OLS regression co-
efficients. The outcome is a dummy = 1 if the
respondent is familiar with the EPA.

Placebo outcomes

Table A40: Repeating the analyses with a placebo institution, the IRS.

	(1)	(2)	(3)	(4)
	IRS improves lives	IRS improves lives	IRS indepen- dence scale	IRS indepen- dence scale
T1: Low search costs & high value info	0.120 (0.0808)	0.0710 (0.0789)	0.183* (0.0946)	0.113 (0.0923)
T2: High search costs & high value info	0.0269 (0.0802)	0.0201 (0.0781)	0.101 (0.0934)	0.0998 (0.0917)
Constant	3.515*** (0.0574)	3.835*** (0.235)	3.117*** (0.0667)	3.941*** (0.263)
Observations	1,628	1,628	1,606	1,606
R-squared	0.002	0.083	0.004	0.082
Demographics	No	Yes	No	Yes
News	No	Yes	No	Yes
Mean dep. var.	3.560	3.560	3.197	3.197

Notes: Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

OLS regression coefficients. In cols (1-2), the outcome reports whether the IRS recommends policies and regulations that improve or worsen the life of the respondents and their loved ones on a 7-point likert scale from "Always worsen" to "Always improve"; in cols (3-4), whether the IRS is politicized versus independent on a 7-point likert scale from "Completely politicized" to "Completely independent". The control variables are reported in the Appendix.

Alternative specification

Table A41: Repeating the analyses with alternative specifications (probit and ordered probit) does not affect results.

	(1)	(2)	(3)	(4)	(5)	(6)
	Knowledge accuracy	Knowledge accuracy	EPA improves lives	EPA improves lives	EPA indepen- dence scale	EPA indepen- dence scale
T1: Low search costs & high value info	0.364*** (0.0864)	0.352*** (0.0879)	0.107* (0.0628)	0.0448 (0.0640)	0.171*** (0.0634)	0.150** (0.0647)
T2: High search costs & high value info	-0.555*** (0.0834)	-0.557*** (0.0854)	0.0279 (0.0622)	0.00432 (0.0625)	0.118* (0.0613)	0.126** (0.0617)
Constant cut1			-2.010*** (0.0767)	-1.909*** (0.187)	-1.228*** (0.0544)	-1.731*** (0.198)
Constant cut2			-1.529*** (0.0600)	-1.392*** (0.185)	-0.495*** (0.0479)	-0.959*** (0.194)
Constant cut3			-0.962*** (0.0507)	-0.790*** (0.182)	0.242*** (0.0472)	-0.188 (0.193)
Constant cut4			-0.259*** (0.0471)	-0.0322 (0.181)	0.903*** (0.0501)	0.500*** (0.192)
Constant cut5			0.388*** (0.0476)	0.678*** (0.182)	1.503*** (0.0579)	1.127*** (0.196)
Constant cut6			1.502*** (0.0596)	1.888*** (0.187)	2.267*** (0.0877)	1.921*** (0.199)
Constant	-0.0287 (0.0596)	-0.211 (0.266)				
Observations	1,400	1,400	1,626	1,626	1,613	1,613
Demographics	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.441	0.441	4.788	4.788	3.360	3.360

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In cols (1-2), the outcome is whether the respondent correctly answered a question asking about the AQI score in their zipcode or the location of the closest EPA office (according to the treatment status); in cols (3-4), whether the EPA recommends policies and regulations that improve or worsen the life of the respondents and their loved ones on a 7-point likert scale from "Always worsen" to "Always improve"; in cols (5-6), whether the EPA is politicized versus independent on a 7-point likert scale from "Completely politicized" to "Completely independent". The control variables are reported in the Appendix.

Heterogeneous effects by EPA's office location

Table A42: Residing in States that do not host an EPA office does not affect respondents' attitudes

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
	Knowledge accuracy	Knowledge accuracy	EPA improves lives	EPA improves lives	EPA independence scale	EPA independence scale	Support for EPA's involvement	Support for EPA's involvement
EPA office same State	0.0370 (0.0322)	0.0323 (0.0346)	0.0659 (0.0843)	-0.0190 (0.0878)	0.0523 (0.0880)	0.00473 (0.0898)	0.246 (1.713)	-1.477 (1.774)
T3 - Unif. treat.	0.0266 (0.0378)	0.0206 (0.0375)	-0.118 (0.0933)	-0.0828 (0.0886)	-0.214** (0.0969)	-0.198** (0.0946)	-3.383* (1.984)	-2.761 (1.897)
T3*EPA office	0.0757 (0.0571)	0.0878 (0.0569)	0.0820 (0.143)	0.119 (0.135)	0.0359 (0.147)	0.0383 (0.144)	2.853 (2.998)	3.488 (2.848)
Constant	0.415*** (0.0224)	0.373*** (0.0972)	4.818*** (0.0573)	4.419*** (0.234)	3.419*** (0.0618)	4.025*** (0.257)	61.45*** (1.205)	61.46*** (4.776)
Observations	1,400	1,400	1,626	1,626	1,613	1,613	1,580	1,580
R-squared	0.017	0.062	0.018	0.154	0.007	0.087	0.003	0.137
Demographics	No	Yes	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.441	0.441	4.788	4.788	3.360	3.360	60.61	60.61

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. In this table, we evaluate whether the treatment T3 has heterogeneous effects across locations with the closest EPA office. In cols (1-2), the outcome is whether the respondent correctly answered a question asking about the AQI score in their zipcode or the location of the closest EPA office (according to the treatment status); in cols (3-4), the outcome reports whether the EPA recommends policies and regulations that improve or worsen the life of the respondents and their loved ones on a 7-point likert scale from "Always worsen" to "Always improve"; in cols (5-6), the outcome reports whether the EPA is politicized versus independent on a 7-point likert scale from "Completely politicized" to "Completely independent". The control variables are reported in the Appendix.

Table A43: Heterogeneous analyses by air pollution at the Zip Code level.

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
	Knowledge accuracy	Knowledge accuracy	EPA improves lives	EPA improves lives	EPA independence scale	EPA independence scale	Support for EPA's involvement	Support for EPA's involvement
T1 - Inform. treat	0.302*** (0.0430)	0.281*** (0.0440)	0.0364 (0.112)	-0.0465 (0.108)	0.0415 (0.121)	-0.0124 (0.118)	1.685 (2.273)	0.476 (2.163)
T1 *AQI>50	0.0993 (0.0609)	0.121** (0.0612)	0.151 (0.169)	0.169 (0.160)	0.0818 (0.176)	0.0769 (0.169)	0.709 (3.446)	0.362 (3.249)
AQI>50	-0.0779** (0.0383)	-0.0797** (0.0396)	-0.177 (0.119)	-0.189* (0.114)	-0.0804 (0.121)	-0.0827 (0.117)	-2.382 (2.422)	-2.135 (2.323)
Constant	0.316*** (0.0276)	0.187* (0.113)	4.875*** (0.0805)	4.368*** (0.324)	3.441*** (0.0858)	3.843*** (0.343)	61.74*** (1.604)	60.68*** (6.092)
Observations	949	949	1,068	1,068	1,060	1,060	1,036	1,036
R-squared	0.131	0.174	0.016	0.158	0.003	0.101	0.004	0.152
Demographics	No	Yes	No	Yes	No	Yes	No	Yes
News	No	Yes	No	Yes	No	Yes	No	Yes
Mean dep. var.	0.423	0.423	4.768	4.768	3.388	3.388	60.34	60.34

Notes: Standard errors in parentheses. *** p<0.01, ** p<0.05, * p<0.1.

OLS regression coefficients. The sample includes only respondents who were asked about the air quality (T1 and T2). In cols (1-2), the outcome is whether the respondent correctly answered a question asking about the AQI score in their zipcode or the location of the closest EPA office (according to the treatment status); in cols (3-4), the outcome reports whether the EPA recommends policies and regulations that improve or worsen the life of the respondents and their loved ones on a 7-point likert scale from "Always worsen" to "Always improve"; in cols (5-6), the outcome reports whether the EPA is politicized versus independent on a 7-point likert scale from "Completely politicized" to "Completely independent". The control variables are reported in the Appendix.